

**REDESIGN OF DATABASE ALGORITHMS  
FOR NEXT GENERATION  
NON-VOLATILE MEMORY TECHNOLOGY**

**Weiwei Hu**

**Bachelor of Engineering  
Tsinghua University, China**

**A THESIS SUBMITTED  
FOR THE DEGREE OF MASTER OF SCIENCE  
SCHOOL OF COMPUTING  
NATIONAL UNIVERSITY OF SINGAPORE**

**2012**

# Acknowledgement

I would like to express my sincere gratitude to my supervisor, Prof. Beng Chin Ooi for his guidance and encouragement throughout the research work. His continuous support led me to the right way. He teaches me the right working attitude which I believe will help me a lot along my whole life. His understanding and guidance provides me a good direction of my thesis.

I would also like to thank Prof. Kian-Lee Tan and Guoliang Li, who help me a lot to edit the paper and give me valuable advice for research.

During this work I have collaborated with many colleagues in the database lab, who have given me lots of valuable comments and ideas. I would like to thank all of them. They are Sai Wu, Su Chen, Vo Hoang Tam, Suraj Pathak, Peng Lu, Yanyan Shen etc. I have really enjoyed the pleasant cooperation and work with these smart and motivated people.

Finally, and most importantly, I would like to thank my parents for their encouragement and support. All of these have helped me a lot to complete my graduate studies and this research work.

# Contents

<b>Acknowledgement</b>	<b>i</b>
<b>Summary</b>	<b>v</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Our Contributions . . . . .	3
1.2 Outline of The Thesis . . . . .	4
<b>2 Next-generation Non-volatile Memory Technology</b>	<b>6</b>
2.1 NVM Technology . . . . .	6
2.2 Positions in the Memory Hierarchy . . . . .	9
2.3 Challenges to Algorithms Design . . . . .	10
2.4 Algorithms Design Considerations . . . . .	12
2.5 Summary . . . . .	13
<b>3 Literature Review</b>	<b>15</b>
3.1 Write-optimized Indexing Algorithms . . . . .	15
3.1.1 HDD-based Indexing . . . . .	16
3.1.2 SSD-based Indexing . . . . .	18
3.1.3 WORM-based Indexing . . . . .	20

3.1.4	PCM-based Indexing . . . . .	20
3.2	Query Processing Algorithms . . . . .	21
3.2.1	Adaptive Query Processing . . . . .	21
3.2.2	SSD-based Query Processing . . . . .	24
3.3	PCM-based Main Memory System. . . . .	26
3.4	Prediction Model . . . . .	27
3.5	Summary . . . . .	29
<b>4</b>	<b>Predictive B<sup>+</sup>-Tree</b>	<b>31</b>
4.1	Overview of the B <sup>p</sup> -tree . . . . .	31
4.1.1	Design Principle . . . . .	32
4.1.2	Basic Idea . . . . .	33
4.2	Main Components of B <sup>p</sup> -tree . . . . .	34
4.2.1	DRAM Buffer . . . . .	35
4.2.2	Warm-up Phase . . . . .	36
4.2.3	Update Phase . . . . .	37
4.3	Summary . . . . .	39
<b>5</b>	<b>Predictive Model</b>	<b>41</b>
5.1	Predictive Model for Warm-up . . . . .	41
5.1.1	Predicting the B <sup>p</sup> -tree Skeleton . . . . .	42
5.1.2	B <sup>p</sup> -tree Construction . . . . .	43
5.2	Predictive Model for Updates . . . . .	48
5.2.1	Search . . . . .	48
5.2.2	Deletion . . . . .	49
5.2.3	Insertion . . . . .	50

5.3	Evaluating $B^p$ -tree . . . . .	53
5.4	Summary . . . . .	56
<b>6</b>	<b>Experimental Evaluation</b>	<b>57</b>
6.1	Experimental Setup . . . . .	57
6.1.1	Experimental platform . . . . .	58
6.1.2	Data sets and workloads . . . . .	58
6.1.3	Algorithms compared . . . . .	59
6.2	Results and Analysis . . . . .	60
6.2.1	Insertion . . . . .	60
6.2.2	Update . . . . .	62
6.2.3	Search . . . . .	63
6.2.4	Node Utilization . . . . .	65
6.2.5	Sensitivity to Data Distribution Changes . . . . .	66
6.3	Summary . . . . .	67
<b>7</b>	<b>Conclusion</b>	<b>70</b>
7.1	Conclusion . . . . .	70
7.2	Future Work . . . . .	71

# Summary

In recent years, non-volatile memory like PCM has been considered an attractive alternative to flash memory and DRAM. It has promising features, including non-volatile storage, byte addressability, fast read and write operations, and supports random accesses. Many research scholars are working on designing adaptive systems based on such memory technologies. However, there are also some challenges in designing algorithms for this kind of non-volatile-based memory systems, such as longer write latency and higher energy consumption compared to DRAM. In this thesis, we will talk about our redesign of the indexing technique for traditional database systems for them.

We propose a new *predictive*  $B^+$ -tree index, called the  $B^p$ -tree. which is tailored for database systems that make use of PCM. We know that the relative slow-write is one of the major challenges when designing algorithms for the new systems and thus our trivial target is to avoid the writes as many as possible during the execution of our algorithms. For  $B^+$ -tree, we find that each time a node is full, we need to split the node and write half of the keys on this node to a new place which is the major source of the more writes during the construction of the tree. Our  $B^p$ -tree can reduce the data movements caused by tree node splits and merges that arise from insertions and deletions. This is achieved by pre-allocating space on the memory

for near future data. To ensure the space are allocated where they are needed, we propose a novel predictive model to ascertain future data distribution based on the current data. In addition, as in [6], when keys are inserted into a leaf node, they are packed but need not be in sorted order which can also reduce some writes.

We implemented the  $B^p$ -tree in PostgreSQL and evaluated it in an emulated environment. Since we do not have any PCM product now, we need to simulate the environment in our experiments. We customized the buffer management and calculate the number of writes based on the cache line size. Besides the normal insertion, deletion and search performance, we also did experiments to see how sensitive our  $B^p$ -tree is to the changes of the data distribution. Our experimental results show that the  $B^p$ -tree significantly reduces the number of writes, therefore making it write and energy efficient and suitable for a PCM-like hardware environment. For the future work, besides the indexing technique, we can move on to make the query processing more friendly to the next generation non-volatile memory.

# List of Tables

2.1	Comparison of Memory Technologies . . . . .	8
4.1	Notations . . . . .	35
6.1	Parameters and their value ranges . . . . .	59



# List of Figures

2.1	PCM Technology . . . . .	7
2.2	PCM Usage Proposals . . . . .	10
4.1	$B^p$ -tree architecture . . . . .	35
4.2	An example of a warm-up phase . . . . .	37
4.3	An example for update phase . . . . .	38
5.1	Warmup Algorithm . . . . .	47
5.2	$B^p$ -tree: Search Operation . . . . .	49
5.3	$B^p$ -tree: Deletion Operation . . . . .	50
5.4	$B^p$ -tree: Insertion Operation . . . . .	54
6.1	Insertion performance . . . . .	61
6.2	Update performance . . . . .	64
6.3	Search performance . . . . .	64
6.4	Leaf nodes utilization . . . . .	66
6.5	Sensitivity to Data Distribution Changes . . . . .	68

# Chapter 1

## Introduction

In the established memory technologies hierarchy, DRAMs and flash memories are two major types that are currently in use, but both of them suffer from various shortcomings: DRAMs are volatile and flash memories exhibit limited write endurance and low write speed. In recent years, we found that the emerging next-generation non-volatile memory (NVM) is a promising alternative to the traditional flash memory and DRAM as it offers a combination of some of the best features of both types of traditional memory technologies. In the near future NVM is expected to become a common component of the memory and storage technology hierarchy for PCs and servers [8, 14, 27].

In this thesis, we will research on how to make the emerging next-generation non-volatile memory adaptive to the existing memory system hierarchy. In particular, we will focus on how to make the traditional database systems work efficiently on the NVM-based memory systems. The problem becomes how should the traditional database systems be modified to best make use of the NVM? This thesis is an initial research on this topic and we will present our design for the indexing

technique. In the future, we can continue to work on redesigning many other components in database systems including query processing, buffer management and transaction management etc.

There are some widely pursued NVM technologies: magneto-resistive random access memory (MRAM), ferroelectric random access memories (FeRAM), resistive random access memory (RRAM), spin-transfer torque memory (STT-RAM), and phase change memory (PCM)[17] and in this thesis, we will focus on PCM technology since it is at a more advanced stage of development and our algorithms can be adapted to other similar memory technologies. We know that there are some differences among the different kinds of NVM technologies, but in the remainder of this thesis, we will use PCM and NVM interchangeably for simplicity and mainly focus on the PCM technology.

Like DRAM, PCM is byte addressable and supports random accesses. However, PCM is non-volatile and offers superior density to DRAM and thus provides a much larger capacity within the same budget [27]. Compared to NAND flash, PCM offers better read and write latency, better endurance and lower energy consumption. Based on these features, PCM can be seen as a form of middleware between DRAM and NAND flash, and we can expect it to have a big impact on the memory hierarchy. Due to its attractive attributes, PCM has been considered a feasible device for database systems [27, 6]. In this thesis, we focus on designing indexing techniques in PCM-based hybrid main memory systems, since indexing will greatly influence the efficiency of the traditional database systems.

There are several main challenges in designing new algorithms for PCM. First, though PCM is faster than NAND flash, it is still much slower than DRAM, especially the write function, which greatly affects system performance. Second, the

PCM device consumes more energy because of the phase change of the material. We will elaborate on this point in Chapter 2. Third, compared to DRAM, the lifetime of PCM is shorter, which may limit the usefulness of PCM for commercial systems. However, as mentioned in [27, 6], some measures could be taken to reduce write traffic as a means to extend the overall lifetime. This, however, may require substantial redesigning of the whole database system. In general, the longer access latency and the higher energy consumption are the major factors that affect the performance of PCM-based memory systems.

## 1.1 Our Contributions

In this thesis, we propose the predictive  $B^+$ -tree (called  $B^p$ -tree), an adaptive indexing structure for PCM-based memory. Our main objective is to devise new algorithms to reduce the number of writes without sacrificing the search efficiency. Our  $B^p$ -tree is able to achieve much higher overall system performance than the classical  $B^+$ -tree in the PCM-based systems. To the best of our knowledge, no paper has addressed the issue in such detail and thoroughness. To summarize, we make the following contributions:

1. We first look into the technology details of the next generation non-volatile memory technology and then we conduct a comprehensive literature review about the indexing design in the traditional database systems which can contribute to our redesign. We also showed the algorithms design consideration for the NVM-based systems.
2. We propose a new predictive  $B^+$ -tree index, called the  $B^p$ -tree, which is designed to accommodate the features of the PCM chip to allow it to work

efficiently in PCM-based main memory systems. The  $B^p$ -tree can significantly reduce both number of writes and energy consumption.

3. We implemented the  $B^p$ -tree in the open source database system PostgreSQL<sup>1</sup>, and run it in an emulated environment. Via these experiments, we will show that our new  $B^p$ -tree index significantly outperforms the typical  $B^+$ -tree in terms of insertion and search performance and energy consumption.

## 1.2 Outline of The Thesis

The remainder of the thesis is organized as follows:

- Chapter 2 introduces the next generation non-volatile memory technology. We provides the technical specifications of some NVM technology and do a comparison between the existing memory technologies and NVM. Based on the understanding of the technical specifications, we present our consideration about how to design adaptive algorithms for NVM-based database systems.
- Chapter 3 reviews the existing related work. In this chapter, we did a comprehensive literature review about the write optimized tree indexing and the existing redesigning work of the NVM-based memory systems. Besides, we also reviewed the existing proposals about the query processing techniques which can be referenced in our future work.
- Chapter 4 presents the  $B^p$ -tree. In this chapter, we propose the main design of  $B^p$ -tree. We talked about the major structure of the tree and gave the algorithms about how to insert and search keys on our  $B^p$ -tree.

---

<sup>1</sup><http://www.postgresql.org/>

- Chapter 5 presents a predictive model to perform the predicting of future data distribution. In the previous chapter, we have present that we need a predictive model to predict the future data distribution and in this chapter, we present the details of our predictive model and show how to integrate the predictive model into our  $B^p$ -tree.
- Chapter 6 presents the experimental evaluation. In this chapter, we did various experiments in our simulation environment and showed that our  $B^p$ -tree reduces both execution time and energy consumption.
- Chapter 7 concludes the paper and provide a future research direction. In this chapter, we concluded our work on the indexing technique and gave some direction of the future work. We can continue to work on many other components of database systems including query processing, buffer management and transaction management etc.

## Chapter 2

# Next-generation Non-volatile Memory Technology

Research work on next-generation non-volatile memory technology has grown rapidly in recent years. Worldwide research and development effort have been made on the emerging new memory devices [38]. Before we move on to what we have done for the new design, we need to be clear about what exactly this emerging new memory technology is. In this chapter, we will review the technical specifications about the PCM technology. Besides, we will also talk about the challenges we face to design algorithms for such systems and further our design considerations and targets.

### 2.1 NVM Technology

There are some widely pursued NVM technologies: magneto-resistive random access memory (MRAM), ferroelectric random access memories (FeRAM), resistive random access memory (RRAM), spin-transfer torque memory (STT-RAM), and




Figure 2.1: PCM Technology

phase change memory (PCM). As they have relatively similar features, in this work we focus on PCM since it is at a more advanced stage of development and it can be expected to come out earlier.

Generally, NVM technologies share some features in common. Most NVM chips have comparable read latency than DRAM and rather higher write latency. They have lower energy consumption but have limited endurance. Next we will review the technology details of PCM technology. The physical technology of other NVMs may be different, but in this work, this is not our major focus and we will mainly discuss about the PCM technology. Since they share some of the common features, our design for PCM can be reused for other NVMs.

PCM is a non-volatile memory that exploits the property of some phase change materials. The phase change material is one type of chalcogenide glass, such as  $Ge_2Sb_2Te_5$  (GST) and it can be switched between two states, amorphous and crystalline by the current injection and the heating element. For a large number of times, the difference in resistivity is typically about five orders of magnitude [31], which can be used to represent the two logical states of binary data. As we can see in Figure 2.1[38], crystallizing the phase change material by heating it above the crys-



Table 2.1: Comparison of Memory Technologies

	DRAM	PCM	NAND	HDD
Density	1X	2-4X	4X	N/A
Read Latency	20-50ns	$\sim 50ns$	$\sim 25\mu s$	$\sim 5ms$
Write Latency	20-50ns	$\sim 1\mu s$	$\sim 500\mu s$	$\sim 5ms$
Read Energy	0.8J/GB	1J/GB	1.5J/GB	65J/GB
Write Energy	1.2J/GB	6J/GB	17.5J/GB	65J/GB
Endurance	$\infty$	$10^6 - 10^8$	$10^5 - 10^6$	$\infty$

tallization temperature ( $\sim 300^\circ\text{C}$ ) but below the melting temperature ( $\sim 600^\circ\text{C}$ ) is called the SET operation. The SET operation will turn GST into the crystalline state which corresponds to the logic ‘1’. Then when continuously heated above the melting temperature, GST turns into the amorphous state corresponding to the logic ‘0’, which is called the RESET operation. Writes on phase change material will come down to the states switch, which incurs high operating temperature and further more latency and energy consumption. However, reads on phase change material just need to keep a much lower temperature, which can be faster and more energy saving.

We present a brief comparison on the properties between PCM and other layers in the memory hierarchy, including DRAM, NAND Flash (Solid State Drives, NAND for abbreviation) and HDD (Hard Disk Drives). Table 2.1 summarizes the properties of these different memory technologies, as presented in recent work [6, 27, 5] and the data all corresponds to raw memory chips.

From Table 2.1, we can see that the PCM has promising characteristics. Compared to DRAM, PCM has a density advantage over DRAM which means more memory capacity within a same size chip and further a lower price per capacity. This cheaper price can lead to orders of magnitude of capacity larger within the same budget. Then the read and write performance is also very efficient. We can

see that the read latency of PCM is comparable to that of the DRAM. Although writes are almost an order of magnitude slower than that of DRAM, some techniques like buffer organization or partial writes could be used in algorithms design to reduce the performance gap. For NAND, the write on NAND should be based on pages and even though only small parts of a page are modified, the whole page need to be rewritten, which is called the erase-before-writes problem. NAND suffers from the erase-before-writes problem greatly and this issue caused the slow read and write speed directly compared to DRAM and PCM. Unlike NAND, PCM uses a totally different technology and it does not have the problem of erase-before-writes and thus supports random reads and writes more efficiently. We can see that reads and writes on PCM are orders of magnitude faster than those of NAND and the endurance is also higher than that of NAND. In summary, in most aspects, PCM can be seen as a technology in the middle between DRAM and NAND Flash. Even though PCM has its own shortcomings but we believe that it will have a major role to play in the memory hierarchy, impacting system performance, energy consumption and reliability because of its promising features.

## 2.2 Positions in the Memory Hierarchy

In previous section, we talked about the technical specifications of the PCM technology and we did a comparison between the several commonly used memory technologies. Now in this section, we want to talk about how to best integrate PCM into the existing memory systems, in other words, what is the proper position of PCM in the current memory hierarchy.

In recent years, the computer systems community has already got various re-

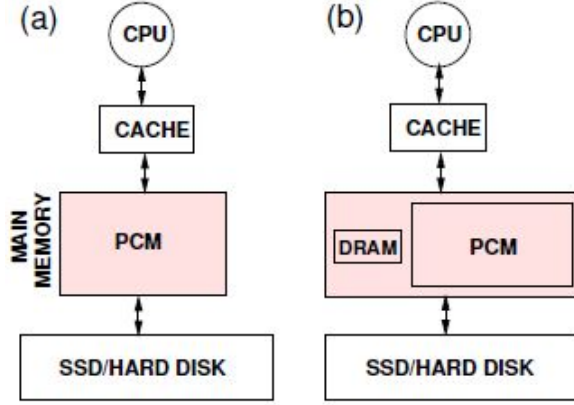


Figure 2.2: PCM Usage Proposals

search proposals on how to make use of PCM technology in the current memory systems. Among all of them, there are mainly two schools of thought [14, 27, 6] as shown in Figure 2.2[6]. One is to replace DRAM with PCM directly to achieve larger main memory capacity [14]. Even though PCM is slower than DRAM, Lee et al. [14] have shown that some optimizations like buffer organization and partial writes can be used to improve the system performance while preserving the high density and non-volatile property. The other proposal is to replace the DRAM with a large PCM and a small DRAM (3% - 8% size of the PCM capacity [27, 30]) and use the DRAM as a buffer to keep some frequently accessed data to improve the system performance. In this work, we will adopt the second approach, with the use of PCM as “extended” but slower main memory and we will use a small number of DRAM as a buffer.

## 2.3 Challenges to Algorithms Design

Now we are clear about the features of PCM. Then before we start to design algorithms, we need to know the major challenges we face and the major target we

want to reach.

From the technologies comparison in previous sections, we can find the following three challenges we need to overcome.

1. Slow writes. Even the read and write speed is much faster than that of the NAND, it is still a bit slower than DRAM which will influence the system efficiency greatly. This challenge is the major one we want to overcome in this thesis. The idea is a bit trivial that since the writes are slow, we want to avoid writes as many as possible.
2. High energy consumption. This challenge is related to the writes. We know that each time we want to write values to the PCM chip, we need to switch its state. Then we need to heat to switch the state of the phase change material which leads to high energy consumption. But for read, since we do not need to switch the state, the energy consumption is much lower.
3. Limited endurance. Existing PCM prototypes have a write endurance ranging from  $10^6$  to  $10^8$  writes per cell [6]. With some good round robin or write leveling algorithms, a PCM main memory can last several years working time [26]. However, such kinds of algorithms should be conducted in the memory driver layer, which is not our main focus then.

From these challenges, we can find that actually if we want to make best use of PCM technology in our existing systems, the most important requirement is to figure out the challenge of high write latency. Our basic idea is that since the speed can not be raised physically, can we just avoid the writes as many as possible?

Then The design objective becomes how to reduce the number of writes in our new algorithms which is an widely studied topic, especially the algorithms designed

for NAND in recent years. However, our consideration is different from that of the algorithms design for NAND. For NAND, they want to avoid the erase-before-writes and thus they will mostly use the batch algorithms to convert random accesses to sequential writes. Our design consideration is different that we can support random writes efficiently but we want to reduce the number of writes including both random writes and sequential write as many as possible. Once the number of writes is limited, we can reduce the energy consumption and extend the life time as well.

## 2.4 Algorithms Design Considerations

Let us go back to our initial problem that we want to make best use of PCM in the existing database systems and we want to integrate PCM into the memory systems. Thus we considering the algorithms design, we need to be careful about the following design goals: (1) CPU time, (2) CPU cache performance and (3) energy consumption. Compared to DRAM, the major shortcoming of PCM is the high write latency. Then for general algorithms, the slow write speed should not influence the cache performance, it can only influence the CPU execution time and energy consumption performance. We also know that the PCM writes incur much higher energy consumption and is much slower than read. Then the major challenge we are facing now is how to reduce the number of writes as many as possible. Actually we have already had this basic direction in mind in previous sections.

Next the problem comes. We need a metric to measure the number of writes on the database algorithms level. In other words, we need to determine what

granularity of writes we need to use in the algorithms analysis using PCM as the main memory. Generally when analyzing algorithms for main memory, we need to consider two granularities including bits and cache lines. For the high level database systems, we have the buffer management and it is easier to count the number of cache line based writes. However, in order to simulate the energy consumption, we need to get the number of writes based on the bits granularity.

Then we can use a mixture of these two metrics. To evaluate the CPU time, we count the number of writes based on the cache line granularity and to evaluate the energy consumption, we first compare the new cache line to the old cache line to get the number of modified bits and get the energy consumption then based on the bits level. Since we have not got any PCM prototype, we need to build a simulated environment to evaluate our algorithms. These should be configured in our simulated environment.

## 2.5 Summary

In this chapter, we introduced the next-generation non-volatile memory technology. There are many kinds of popular non-volatile memory technologies and in this work, we will mainly focus on the phase change memory (PCM), but our algorithms can also be adaptive to other non-volatile memories having the similar features. We present the technical specification details of PCM and did a comparison among PCM and some commonly used memory technologies like DRAM, NAND and HDD about the major features. We found that PCM has its advantages, but there are also some challenges we need to overcome when designing algorithms for PCM-based database systems. Our main design goal is to avoid the slow writes as

many as possible, which further can reduce the energy consumption and extend the lifetime of PCM chips. Finally, we discussed about some metrics to evaluate our new algorithms.

# Chapter 3

## Literature Review

Database research community has contributed a lot to the algorithms design of the traditional database systems in the last several decades. Database system is also very complex and there are many components inside each of which can be worthy of lots of effort to work on. In this work, we mainly focus on the indexing technique and query processing algorithms. In this thesis, we have proposed a new design of  $B^+$ -tree technique and we will leave the query processing to the future work.

In this chapter, we did the literature review, including the write-optimized indexing algorithms, traditional query processing algorithms and some recent PCM-based main memory system proposals. Since our indexing design has a prediction model, we also reviewed some works on histograms and how to use histograms to construct the prediction model.

### 3.1 Write-optimized Indexing Algorithms

In previous chapters we know that our main design goal is to reduce the number of writes, thus we first did a brief survey of the write-optimized indexing techniques.



We want to find out whether these existing techniques can be used to our new PCM-based systems or not and if they can not be used directly, whether we can borrow some ideas from them. Actually write-optimized  $B^+$ -tree index has been an intensive research topic for more than a decade. In this section, we will review the write-optimized indexing algorithms for HDD, SSD technologies and their design consideration is similar to ours, but there are also some differences. For HDD-based indexing, the proposed solutions are mainly focusing on using some DRAM buffer to convert small random writes into sequential writes. For SSD-based indexing, some of the proposed solutions change the major structure of  $B^+$ -tree and some solutions add an inner layer using SSD between DRAM and HDD, but their major design consideration is very similar and they expect to avoid the erase-before-writes as much as possible and they also expect to convert the small random writes into sequential writes in some sense.

### 3.1.1 HDD-based Indexing

For the  $B^+$ -tree index on hard disks, there are many proposals to optimize the efficiency of write operations and logarithmic structures have been widely used. In [21], O’Neil et al. proposed the LSM-tree to maintain a real-time low cost index for the database systems. LSM-tree (Log-Structured Merge-tree) is a disk-based data structure and it was designed to support high update rate over an extended period efficiently. The basic idea of LSM-tree is to use an algorithm to defer and batch the random index changes to reduce the disk arm movements. Some smaller components of LSM-tree will be entirely memory resident and the larger and major components will be disk-based. The smaller components resident in memory will be used as a buffer to keep the frequently referenced page nodes in the

larger components. The insertion to the memory resident component has no I/O cost and the larger component on disk is optimized for sequential disk access with nodes 100% full. This optimization is similar to that used in the SB-tree [22]. Both of them support multi-page reads or writes during a sequential access to any node level below the root, which can offer high-performance sequential disk access for long range retrievals. Each time when the smaller component reaches a threshold size near the maximum allotted, it will be merged into the large components on the disk. A search on the LSM-tree will search both the smaller components in the memory and the larger components on the disk. Then LSM-tree is most useful in applications where index insertion are more than searches, which is just the case for history tables and log files etc.

In [3], Lars Arge proposed the buffer tree for the optimal I/O efficiency. The structure of the buffer tree is very similar to that of the traditional  $B^+$ -tree. The major difference is that in buffer tree, there is a buffer in the main memory for each node on the hard disk. When we want to update the tree, the buffer tree will construct an entry with the inserted key, a time stamp and an indication of whether the key is to be inserted or deleted and put the entry into the buffer of the root. When the main memory buffer of the root is full, it will insert the elements in the buffer downwards to its children and this buffer-emptying process will be done recursively on internal nodes. The main contribution of the buffer tree is that it is a simple structure with efficient I/O operations and can be applied to other related algorithms.

Graefe proposed a new write optimized  $B^+$ -tree index in [10] based on the idea of the log-structured file systems [32]. Their proposals make the page migration more efficient and retain the fine-granularity locking, full concurrency guarantees

and fast lookup performance at the same time.

Most of the HDD-based write optimized tree indexing follow the following two idea. First, most of them want to convert the many random writes to batch sequential writes to raise the efficiency; second, a small area of DRAM can be used as the buffer. For our design, we want to support both random and sequential writes efficiently, thus we cannot “hold” the random writes to a sequential write. But the idea of using DRAM as buffer can be used as well. In previous chapters, we have found out that a small area of DRAM buffer can make our system much more efficient.

### 3.1.2 SSD-based Indexing

Recently, there are some proposals on the write-optimized  $B^+$ -tree index on SSDs [1][15][39]. The major bottleneck of the  $B^+$ -tree index for SSDs to overcome is the rather slower small random writes because of the erase-before-write requirement.

In [39], an efficient  $B$ -tree layer (BFTL) was proposed to handle the fine-grained updates of  $B$ -tree index efficiently. BFTL is introduced as a layer between file systems and FTL and thus there is no need to modify the existing applications. BFTL is considered as a part of the operating system. BFTL consists of a small reservation buffer and a node translation table.  $B^+$ -tree index services call from the upper-level applications are handled and translated from the file system of operation system to BFTL and then a block-based calls are sent from BFTL to FTL to do the operation. When a new record is inserted or updated to the  $B^+$ -tree, it will first be temporarily held by the reservation buffer of BFTL and then flushed in batch operation to reduce the writes latency.

FlashDB was proposed in [20] and it is a self-tuning database system optimized

for sensor networks using flash SSDs. The self-tuning  $B^+$ -tree index in the FlashDB uses two modes, Log and Disk, to make the small random writes together on consecutive pages. The  $B^+$ -tree (Disk) assumes that the storage is a disk-like block-device. The disadvantage is that updates are expensive. Even if only a small part of the node needs to be updated, the whole block needs to be written. Thus  $B^+$ -tree (Disk) is not suitable for write-intensive workload.  $B^+$ -tree (Log) design is a log-structured-like indexing and it can avoid the high update cost of  $B^+$ -tree (Disk). The basic idea is to construct the index tree as transaction logs. Updates will be put into buffer first and flushed into SSD when the buffer contains enough data to fill a page.

More recently, Li et al. [15] proposed the FD-tree which consists of two main parts, a head  $B^+$ -tree in the DRAM and several levels of sorted runs in the SSDs. Thus the basic idea is to limit the random writes to the small top  $B^+$ -tree and then merge into the lower runs after they have been transformed into sequential writes. FD-tree modifies the basic structure of a traditional  $B^+$ -tree and their major contribution is to limit random writes to a small area and further raise the insertion efficiency.

For the SSD-based indexing, some of them tried to use the same idea of that of the HDD-based indexing. We can adopt the idea of using DRAM as a buffer. FD-tree is also efficient in some sense, but it modified the basic structure of the  $B^+$ -tree, which is not what we want. We want that our  $B^+$ -tree can be easily adapted to the existing traditional database systems and thus we do not want to modify the main structure too much.

### 3.1.3 WORM-based Indexing

There are also some proposals focusing on the Write-Once-Read-Many (WORM) storage [16][25]. In [16], Mitra, Hsu and Winslett proposed a novel efficient trustworthy inverted index for keyword-based search and a secure jump index structure for multi-keyword searches. In [25], Pei et al. proposed the TS-Trees and they also built the tree structure based on a probabilistic method. These WORM indexing proposals mainly focused on designing mechanisms to detect adversarial changes to guarantee trustworthy search. Unlike WORM indexing, in PCM-based indexing, we want to reduce the number of writes. Moreover, we can update the index and afford the small penalty of adjustments due to data movement if the prediction is no longer that accurate because of changes in data distributions over time.

### 3.1.4 PCM-based Indexing

The recent study [6] has outlined new database algorithm design considerations for PCM technology and initiated the research on algorithms for PCM-based database systems. To our best knowledge, this paper is the most relevant to our work until now. In the paper, Chen, Gibbons and Nath described the basic characteristics and the potential impact of PCM on the database system design. They presented analytic metrics for PCM endurance, energy and latency, and proposed techniques to modify the current  $B^+$ -tree index and Hash Joins for better efficiency on the PCM-based database system. Their idea is to unsort the keys in the node of  $B^+$ -tree which can reduce large number of writes. In our proposal, we will take a step further in this direction and design a PCM-aware  $B^+$ -tree, called the  $B^p$ -tree and will compare our  $B^p$ -tree with their  $B^+$ -tree.

## 3.2 Query Processing Algorithms

Query processing is an important part of the traditional database systems. The aim of query processing is to transform a query in a high-level declarative language (e.g. SQL) into a correct and efficient execution strategy. Different execution strategies can lead to much difference in execution efficiency. The traditional query processing algorithms include two types, one is heuristic-based query optimization and the other is cost-based query optimization. In heuristic-based query optimization, given a query expression, the algorithm will perform selections and projections as early as possible and it will try to eliminate duplicate computations. In cost-based query optimization, the algorithm will estimate the cost of different equivalent query calculator and choose the execution plan with the lowest cost estimation.

In this section, we will review the existing query processing algorithms. We mainly focus on two parts, the adaptive query processing and the recently proposed query processing for SSD-based database systems. This section of review will give us some inspiration about what to do in the future and we will discuss about it in detail in Chapter 7.

### 3.2.1 Adaptive Query Processing

In traditional query processing, the query optimizer can not have necessary statistics during the compile time, thus it may lead to poor performance, especially in long running query evaluations. Adaptive query processing addresses this problem and the idea is to adapt the query plan to changing the environmental conditions at runtime. In [12], adaptive query processing is defined as that if the query processing system receives information from its environment and determines its behavior

according to the information in an iterative manner, that means there is a feedback loop between the environment and the behavior of the query processing system. The research on adaptive query processing mainly lies on two main directions. One is to modify the execution plan at runtime according to the changes of the evaluation environment, the other is to develop new operators that has more flexibility to deal with unpredictable conditions. Then we will review some of the classical proposals on adaptive query processing.

### **Memory Adaptive Sorting and Hash Join**

Memory shortage is a common design restriction for query processing techniques, especially for sorting and join since they need large amount of excess memory. In [24], Pang et al. introduces new techniques for external sorting to adapt to fluctuations in memory availability. Since memory buffer size can greatly influence the performance of sorting, memory-friendly management strategies need to be taken. [24] introduces a dynamic splitting technique, which adjusts the buffer size to reduce the performance penalty due to the memory shortages. The basic idea is to split the merge step of sorting run into some smaller sub-steps in case of memory shortages and when the memory buffer is larger, it will combine some small sub-steps into larger steps. This adjust is adaptive and is balancing well. For hash join, [23] proposes partially preemptible hash joins (PPHJs) which is one kind of memory adaptive hash joins. The idea is similar to that of [24], they split the original relations and if the memory buffer is not enough, it will flush part of the partition to disk. The most efficient case for PPHJs is when the inner relation can be put in memory but the outer relation can only be scanned and partially put into the buffer. It can reduce both I/O and the total response time. [40] is also a

memory-adaptive sorting which is complementary to [24]. It allows many sorts to run concurrently to improve the throughput, while [24] focuses on improving the query response time.

### **Operators for Producing Partial Results Quickly**

In most database applications, we focus on the total response time of all the results. But in some applications especially in online aggregation, it is important to get some of the results in a very short time and respond earlier while leaving the remaining process running at the same time. To this end, pipelining algorithms are used for the implementation of join and sort operators. Ripple joins [11] are a new family of physical pipelining join operators. They make use of both block nested loops join and hash joins. They target on online processing of multi-table aggregation queries in traditional DBMS. It is designed to minimize the time until an acceptably precise estimate of the query result is available, as measured by the length of a confidence interval. Ripple join comes from the nested loops join since it has an outer relation and an inner relation, the difference is that it adjusts the retrieve rates of tuples from these two input based on the statistical information during the runtime and reduce the response time of important part of the whole results. Xjoin [37] is a variant of Ripple joins. It partitions the input and thus requires less external memory, which makes it more suitable for parallel processing. [29] proposes another pipelining reorder operator for providing user control during long running, data intensive operations to get partial results during the process. The input of this operator is an unordered set of data and it can produce a nearly sorted result according to user preferences which can change during the runtime with an attempt to ensure that interesting items are processed first.



### **Algorithms that Defer Some Optimization Decisions until Runtime**

Sometimes the statistical information gathered in the beginning of the query processing is not that accurate which can lead to a bad performance. There are some algorithms that defer some optimization decisions until they have collected enough statistics. [13] proposes an algorithm that detects sub-optimality in query plans at runtime, through on-the-fly collection of query statistics. It can improve the total performance by either reallocating resources (e.g. memory) or by reordering the query plan. It introduces a new operator called statistics collector operator which is used for the re-optimization algorithm. The algorithm is heuristics-based and relies heavily on intermediate data materialization. The new operator is inserted into the query plan, ensuring that it does not slow down the query by more than a specific fraction and also assigns a potential inaccuracy level of low, medium or high to the various estimates, which will be used in the following processing.

### **3.2.2 SSD-based Query Processing**

Recently there are some research work on query processing for SSD-based database systems. In previous Section 3.1.2, we have said that the major design goal of SSD-related algorithms is to reduce the influence of the high write latency caused by the erase-before-write restriction. We will then review the following several proposals about query processing algorithms.

In [36], Graefe et al. focus on the impact of SSD characteristics on query processing in relational databases and especially on join processing. They first demonstrate a column-based layout within each page and show that it can reduce the amount of data read during selections and projections. Then they introduce

FlashJoin, which is a general pipelined join algorithm that minimized accesses to base and intermediate relational data. FlashJoin has better performance than a variety of existing binary joins, mainly due to its novel combination of three well-known ideas: using a column-based layout when possible, creating a temporary join index and using late materialization to retrieve the non-join attributes from the fewest possible rows. Then we focus on the details of FlashJoin algorithm. FlashJoin is a multi-way equi-join algorithm, implemented as a pipeline of stylized binary joins, each of which includes a join kernel and a fetch kernel. The join kernel computes the join and outputs a join index, which is used in the fetch kernel to do a late materialization, which only retrieves the needed attributes to compute the next join using RIDs specified in the previous join index. The final fetch kernel retrieves the remaining attributes for the result. The experiments show that FlashJoin significantly reduces memory and I/O requirements for each join in the query and raises the query performance greatly.

[35] propose RARE-join algorithm. They convert traditional sequential I/O algorithms to ones that use a mixture of sequential and random I/O to process less data in less time. To make scans and projection faster, they examine a PAX-based page layout [2], which arranges rows within a page in column-major order. Then they designed the RARE-join (RANdom Read Efficient Join) algorithm for the column-based page layout. RARE-join first constructs a join index and then retrieves only the pages and columns needed for computing the join result. The main idea of RARE-join is very similar to that of the FlashJoin. In [4], the authors show that many of the results from magnetic HDD-based join methods also hold for flash SSDs. Their results show that in many cases the block nested loops join over sort-merge join and grace hash join works well for SSD and they also propose

the idea that simply looking at the I/O costs when designing new flash SSD join algorithms can be problematic, because the CPU cost can be a big part of the total join cost in some cases. Their idea also tells us that we need to do a precise research on the new algorithms design and we need to consider all kinds of costs which can influence the main system performance in different aspects.

### 3.3 PCM-based Main Memory System.

Several recent studies from the computer architecture community have proposed new memory system designs on PCM. They mainly focused on how to make PCM a replacement or an addition to the DRAM in the main memory system. Although these studies mainly focused on the hardware design, they provided us the motivation on the use of PCM in the new memory hierarchy design for database applications.

The major disadvantages of the PCM for a main memory system are the limited PCM endurance, longer access latency and higher dynamic power compared to the DRAM. There are many relevant studies addressing these problems [27, 41, 14, 26]. In [27], Qureshi, Srinivasan and Rivers designed a PCM-based hybrid main memory system consisting of the PCM storage coupled with a small DRAM buffer. Such an architecture has both the latency benefits of DRAM and the capacity benefits of PCM. The techniques of partial writes, row shifting and segment swapping for wear leveling to further extend the lifetime of PCM-based systems have been proposed to reduce redundant bit-writes [41, 14]. Qureshi et al. [26] proposed the Start-Gap wear-leveling technique and analyzed the security vulnerabilities caused by the limited write endurance problems. Their proposal requires less than 8 bytes of

total storage overhead and increased the achievable lifetime of the baseline system from 5% to 50% of the theoretical maximum. Zhou et al. [41] also focused on the energy efficiency and their results indicated that it is feasible to use PCM technology in place of DRAM in the main memory for better energy efficiency. There are other PCM related studies such as, [33] focusing on error corrections, and [34] focusing on malicious wear-outs and durability.

Many of the studies in computer architecture community focus on the lifetime issue of PCM technology which is in a lower level consideration. This is the reason why we do not focus on the wear out problem in our proposal. This issue should be addressed in the PCM driver level and what we want to do is about the upper level database system algorithms.

### 3.4 Prediction Model

In this section, we will review the traditional prediction model based on the histogram since we need an accurate prediction model in our proposal for indexing algorithms.

Prediction has always been an extremely important activity and it is an important research topic in statistical theory and it is a big business. The basic idea of prediction is to predict the future value or distribution of a random variable (RV). The probability distributions of many RVs encountered in practice are subject to changes over time and thus it is well known that the fundamental goal of predicting is actually to update the probability distribution based on the present and past information.

Traditionally, we use histogram to represent the distribution of a random vari-

able. We divide the observed range of variation of the RV into a number of value intervals and the relative frequency of each interval is defined to be the proportion of observed values of the RV that lie in that interval. We define the relative frequency in each value interval  $I_i$  is the estimate of the probability  $p_i$  that the RV lies in that interval.

Let  $I_1, \dots, I_n$  be the value intervals with  $u_1, \dots, u_n$  as their midpoints and  $p = (p_1, \dots, p_n)$  is the probability vector of the distribution of RV. Then we can get  $\bar{\mu}$  and  $\bar{\sigma}$  as the estimates of expected value  $\mu$  and standard deviation of the RV.

$$\bar{\mu} = \sum_{i=1}^n u_i p_i, \bar{\sigma} = \sqrt{\sum_{i=1}^n p_i (u_i - \bar{\mu})^2} \quad (3.1)$$

Then we can define the most commonly used probability distribution in decision making, the normal distribution. It is completely specified by the above two parameters. It is symmetric around the mean and the probabilities corresponding to the intervals  $[\mu - \sigma, \mu + \sigma]$ ,  $[\mu - 2\sigma, \mu + 2\sigma]$ ,  $[\mu - 3\sigma, \mu + 3\sigma]$  are 0.68, 0.95, 0.997 respectively.

In our prediction model in practice, the probability distributions of RVs may change with time. For example, the distribution of the values inserted into the  $B^+$ -tree may change with time. We want to capture all the dynamic changes occurring in the shapes of probability distributions from time to time and the traditional method is not adequate anymore.

In [19], Murty proposes to represent probability distribution by the traditional distributions to make all changes possible. When updating the distribution of RV, we can change the values of any  $p_i$  which makes it possible to capture any change in the shape of the distribution. In their model, in addition to the present distri-

bution vector  $p_1, \dots, p_n$ , they add the recent histogram  $f_1, \dots, f_n$  and the updated distribution  $x_1, \dots, x_n$ .

$f = (f_1, \dots, f_n)$  represents the estimate of the probability vector in the recent histogram, but it is based on the most recent few observations (for example 50).  $p = (p_1, \dots, p_n)$  is the probability vector in the traditional distribution at the previous updating.  $x = (x_1, \dots, x_n)$  is the updated probability vector which can be obtained by incorporating the changing trend reflected in  $f$  into  $p$ . In [18], the authors proposed to compute  $x$  from  $p$  and  $f$  using the following formula.

$$x = \beta p + (1 - \beta)f \quad (3.2)$$

In equation 3.2,  $\beta$  is a weight between 0 and 1 and they found that  $\beta=0.8$  or 0.9 works well for the model.

We can find that the basic idea of Murty's proposal is to combine the latest change trend of the data distribution into the current distribution to better predict the future distribution. Since they will keep updating the trend vector, the prediction will keep changing based on the current distribution changing trend which makes the prediction more accurate. Their work may be helpful to our prediction model to predict the future inserted values into the  $B^+$ -tree.

### 3.5 Summary

In this chapter, we did a comprehensive literature review on the related topics to our problem. Since we want to design a write-optimized indexing technique, we first reviewed some of the existing write-optimized indexing techniques for other memory technologies based database systems. Then we reviewed the typical query

processing algorithms which can be useful for our future work. After that, we talked about the recent research proposals on the PCM-based memory systems design. We have the similar challenges but we are doing the job on a different level and thus our major focus is a bit different. But some of the basic ideas coming from the computer architecture community can also be used by our design. Lastly, we did a brief review about prediction model since we need to use a prediction model to predict the future database distributions in our design.

# Chapter 4

## Predictive $B^+$ -Tree

In previous chapters, we introduced the background information about phase change memory technology and the algorithms design consideration for PCM-based database systems. From this chapter, we are going to present the design of our predictive  $B^+$ -tree, the  $B^p$ -tree. The purpose of our  $B^p$ -tree is to reduce the number of writes of traditional  $B^+$ -tree while keep the insert and search performance at the same time. Our basic idea is to reduce the number of node splits caused by being full since node splits are a major source of the unnecessary writes of keys. As there are many  $B^+$ -tree variants in database research community, for simplicity, in our work, we will use the standard  $B^+$ -tree and our techniques can be easily extended to other variants.

### 4.1 Overview of the $B^p$ -tree

In this section, we will introduce the basic concept of  $B^p$ -tree. We first talk about the design principle of  $B^p$ -tree and then we present the basic idea of  $B^p$ -tree. The idea itself is very simple but we need to be very careful to make the tree balance



and stable.

### 4.1.1 Design Principle

Our goal is to reduce the number of writes for data insertions and updates, without sacrificing the performance of search queries. We have seen that there are many traditional methods to do write optimization and we can borrow the DRAM buffer idea. For our own design, we want to reduce the number of writes in the following two ways. First, we adopt the **Unsorted Leaf** strategy in [6]. Essentially, newly inserted keys are simply appended to the end of the key entries. As such, they are not necessarily in sorted order which may reduce large amount of writes. Hence, the search cost may incur additional overhead as all entries within a leaf node have to be examined. But since we now consider the main memory algorithms design and we want to set the size of the node to several cache lines' size, this additional overhead will not be that much. Second, we develop a scheme that minimizes data movement caused by node splits and merges. We have present in the previous sections that the splits and merges are the major source of many additional writes required. If we can reduce the number of splits, the performance can be greatly raised. But in the traditional  $B^+$ -tree, the algorithm is very stable that only when the node is full, the split happens. Or only when the node becomes “underflow” because of deletion, we merge the sibling nodes. Now we want to reduce these operations, we need to “break” some of the existing rules. For the remainder of this thesis, we shall focus on reducing the number of node splits and merges, which is the main motivation for designing the  $B^p$ -tree and we will talk about more details next.

### 4.1.2 Basic Idea

The general idea is to predict the data distribution based on the past insertions and pre-allocate space on PCM for accommodating future tree nodes, which can reduce the key movements caused by node splits and merges. At the same time, it is possible to fail some of the basic balance properties of  $B^+$ -tree, but we will adopt some strategies to ensure the balance property and make sure that the difference will not influence the performance. Figure 4.1 illustrates the main architecture of a  $B^p$ -tree. We use the following techniques to implement a  $B^p$ -tree.

#### DRAM Buffer

We use a small DRAM buffer to maintain a small  $B^+$ -tree for current insertions. We also record the summary of previously inserted keys in a histogram and use them to predict the structure of the  $B^p$ -tree. If the buffer is full, we will merge it into the  $B^p$ -tree on PCM.

#### $B^p$ -tree on PCM

Like a standard  $B^+$ -tree, a  $B^p$ -tree is also a balanced multiway search tree. The key differences between the  $B^p$ -tree and the  $B^+$ -tree include:

1. The structures and nodes in a  $B^p$ -tree can be pre-allocated.
2. Given a branching factor  $2M$  of a  $B^p$ -tree, the number of children of an internal node may be smaller than  $M$ . and the real number of children is between  $[0, 2M]$ .
3. The insertions and deletions are different from the  $B^+$ -tree (see Section 5.2).

4. The tree construction process consists of two phases: (i) **warm-up phase**: The first  $N$  keys are initially inserted into the tree as a warm-up process; (ii) **update phase**: All new keys are first inserted into a DRAM buffer. Each time the buffer is full, the keys in DRAM would be merged into the main tree on PCM. For a search query, we will find them from both the  $B^+$ -tree in DRAM and the  $B^p$ -tree in PCM.

## Predictive Model

Predictive model is a very important component in our  $B^p$ -tree. Each time we want to merge the small  $B^+$ -tree on DRAM to PCM, we need to use the predictive model. If the model is too sparse, it may lead to many more nodes and the nodes utilization is very low. If the model is too strict, the influence of the our strategy may be not that obvious and our  $B^p$ -tree becomes similar to the traditional  $B^+$ -tree. Thus an accurate predictive model is very important to our design and we also proposed some strategies to guard the prediction model in order to make it work properly. Currently in our predictive model, we use the histogram and we can get better predictive model for better performance in the future.

## 4.2 Main Components of $B^p$ -tree

In this section, we will describe the details of the construction process of a  $B^p$ -tree. It consists of two phases, namely the warm-up phase and update phase, which will be described in Section 4.2.2 and Section 4.2.3 respectively.

For ease of presentation, we summarize the notations used throughout this paper in Table 4.1.

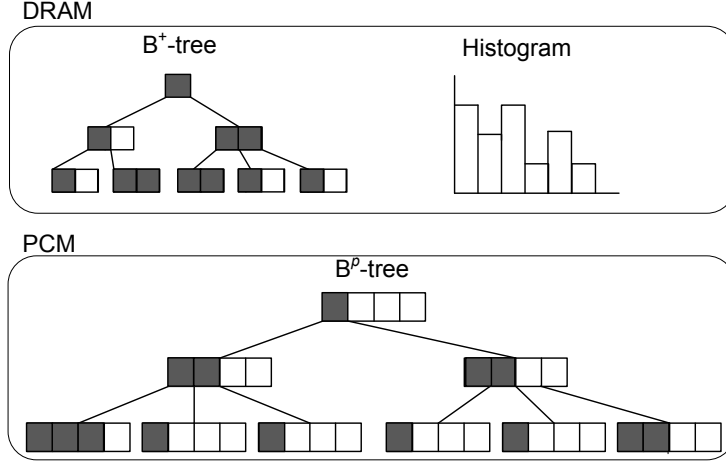


Figure 4.1:  $B^p$ -tree architecture

Table 4.1: Notations

Parameter	Description
$h$	Height of the $B^p$ -tree and $B^+$ -tree
$2M$	The branching factor of the $B^p$ -tree on PCM
$2m$	The branching factor of the $B^+$ -tree on DRAM
$K$	$M$ divided by $m$ ( $K$ is an integer and $K \geq 1$ )
$B_i$	The $i$ -th bucket
$n_i$	Number of entries in the $i$ -th bucket

### 4.2.1 DRAM Buffer

As new keys are inserted into the the DRAM buffer continuously, a small standard  $B^+$ -tree with branching factor  $2m$  is built in the DRAM buffer. If the buffer is full, we will flush the keys in the  $B^+$ -tree to the  $B^p$ -tree on the PCM.

To capture the data distribution, we also maintain a histogram. Suppose the range of the keys is  $[L, U]$ . If we want to partition the keys into buckets  $B_1, B_2, \dots, B_{|B|}$ , the bucket width is  $\frac{U-L}{|B|}$ . For each bucket  $B_i$ , we maintain the number of keys that fall in this bucket, denoted by  $n_i$ . We will use the histogram

to “forecast” the data distribution (Section 5.1).

The main function of DRAM buffer is to adaptively adjust our predictive model based on the currently inserted keys in a time window. Then we can use the updated predictive model to merge all the keys in the time window in the  $B^+$ -tree to the  $B^p$ -tree on PCM.

### 4.2.2 Warm-up Phase

Initially, the  $B^p$ -tree on PCM is empty. We use a DRAM buffer for warm-up. We create a standard  $B^+$ -tree for supporting insertions, deletions and search. Before the buffer is full, we use the conventional  $B^+$ -tree for the initial operations. For the first time that the DRAM buffer is full, all the keys in the buffer will be moved to the PCM, and this step is called the warm-up process. The main function of the warm-up phase is to construct the skeleton of the  $B^p$ -tree on PCM.

Suppose the DRAM buffer can accommodate  $N$  keys. We first predict the total number of possible keys. Then, for each  $B^+$ -tree node, we use our predictive model to decide whether to split it in an eager manner to avoid writes for subsequent insertions. We will provide the details for constructing the initial  $B^p$ -tree in Section 5.1.

Figure 4.2 shows an example for the warm-up phase. The  $B^+$ -tree and histogram are in the DRAM and  $B^p$ -tree is in the PCM. In this example,  $N$  is 10 and the buffer is full and thus we need to flush all the keys in the  $B^+$ -tree to the PCM. The black portion of the histogram bar indicates the number of inserted keys in each range so far, while the whole bar indicates the predicted number of keys in each range based on our predictive model. From this figure, we can observe that the structure of the  $B^p$ -tree is similar to that of the original  $B^+$ -tree. However, there

### Warm-up

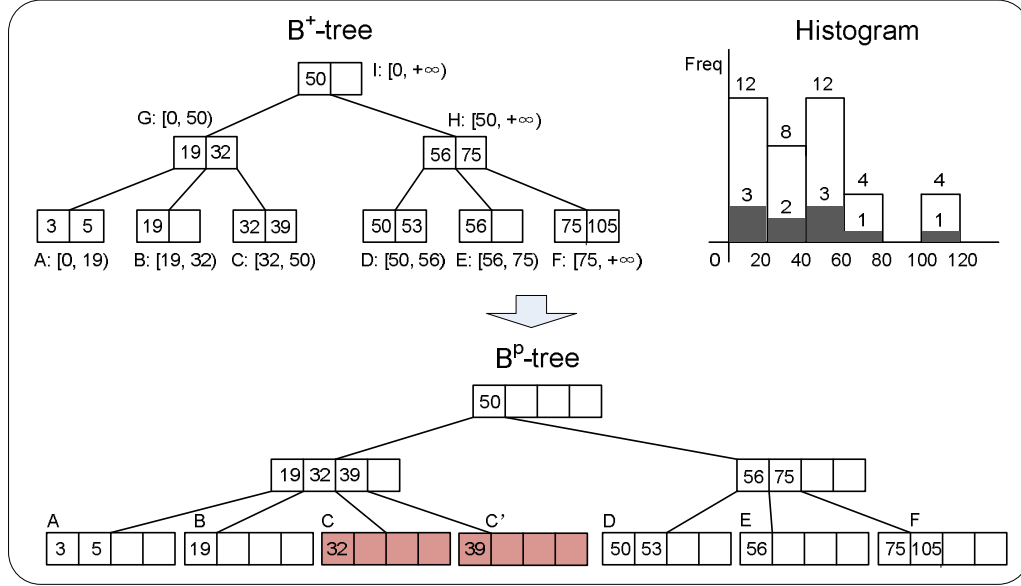


Figure 4.2: An example of a warm-up phase

are two key distinctions. First, the node could be split in an early manner if it meets the requirement of node splits. Second, some of the nodes could *underflow* due to either an enlargement of the node size or an early split. These are guided by our predictive model and tree construction strategy. In the example, node  $C$  in the  $B^+$ -tree is split into node  $C$  and node  $C'$  when it is moved to the  $B^p$ -tree, nodes  $B$  and  $E$  underflow because of the enlargement of the node size, while node  $C$  and node  $C'$  underflow because of the early split. Details about the early split algorithm will be presented in Section 5.1.

### 4.2.3 Update Phase

After the warm-up phase, we have a  $B^p$ -tree structure on the PCM. Then for new operations, we use both the DRAM buffer and  $B^p$ -tree to handle the operations. For an insertion, we insert it into the  $B^+$ -tree. For a search query, we search the key

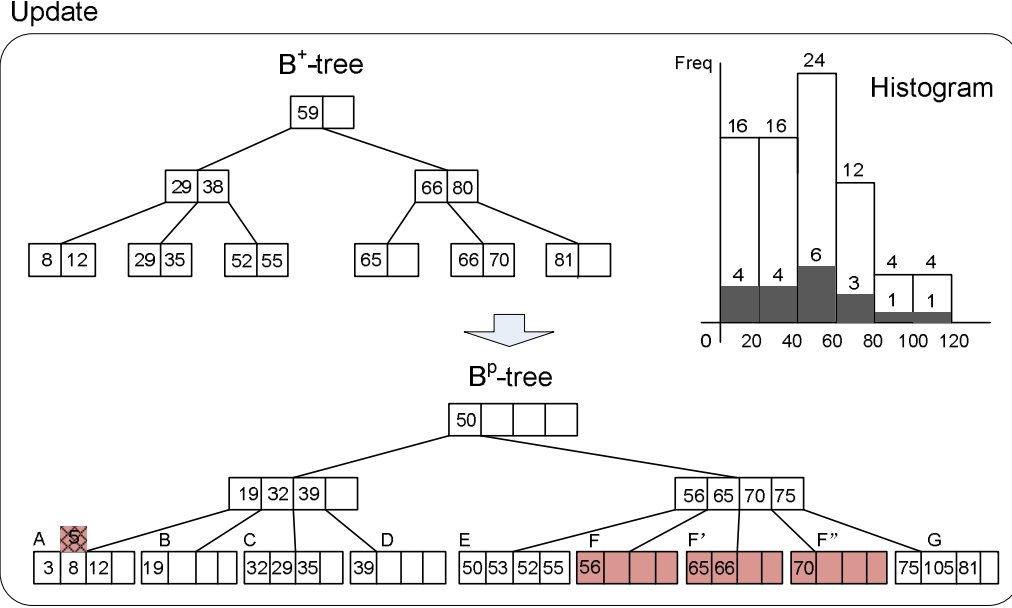


Figure 4.3: An example for update phase

from both the  $B^+$ -tree on DRAM and the  $B^p$ -tree on PCM. If we find it, we return the answer; otherwise we return “null”. (Section 5.2.1). For delete, we search it from both the  $B^+$ -tree and the  $B^p$ -tree. If we find it, we remove it from the  $B^+$ -tree and the  $B^p$ -tree (Section 5.2.2). However, even if a node “underflows” after deletions, we do not merge it with its siblings. The reason is that since the read latency of PCM is much less than the write latency, the overhead caused by empty nodes during query processing is negligible. Furthermore, space could be reserved for future insertion keys to reduce subsequent writes. For update operation, like other indexes, we treat it as a deletion operation followed by an insertion. The deletion operation does not need to be buffered, while the following insertion needs to be buffered first like the standard insertion operation on the  $B^p$ -tree. Note that we need to update the histogram for the insertion and deletion operations. If the DRAM buffer is full, we need to merge the  $B^+$ -tree into the  $B^p$ -tree (Section 5.2.3).

Figure 4.3 shows an example for update phase affected on the earlier example

described in Figure 4.2. The case in Figure 4.3 is that the buffer is full for the second time and all the keys in the  $B^+$ -tree are merged into the  $B^p$ -tree described in Figure 4.2. In this example for the update phase, we want to delete the key 5 in the  $B^p$ -tree index from Figure 4.2. First, we search the  $B^+$ -tree in the buffer and cannot find it. Then we search the  $B^p$ -tree on the PCM and find it in node  $A$  and subsequently remove it from the  $B^p$ -tree. As can be seen from the figure, the histogram is updated to reflect the effect of this deletion and the new round of prediction is performed based on all the keys inserted currently including the keys in the buffer. Node  $F$  in the  $B^p$ -tree is split because of the similar reason as that of the node  $C$  in Figure 4.2. We will describe the details in Section 5.2.

### 4.3 Summary

In this chapter, we introduced our  $B^p$ -tree. We present the design principle and basic idea of  $B^p$ -tree. We know that there are two major parts of  $B^p$ -tree including the DRAM buffer and the normal  $B^+$ -tree on PCM. When new keys are inserted into the  $B^p$ -tree, there are two phases. First, the new key is inserted into the small  $B^+$ -tree on the DRAM buffer, when the DRAM buffer is full for the first time, we merge the whole tree to PCM and construct the skeleton of the tree on PCM based on the predictive model. After that each time when the DRAM is full, we will merge the tree to the existing  $B^+$ -tree on PCM. During the construction, the node of our  $B^p$ -tree may be “underflow” which is different from the traditional  $B^+$ -tree. The reason is that sometimes we want to split the node in advance based on the prediction model even the node is not full. In our design, even some nodes may be “underflow”, we will still use some strategy to ensure that the whole tree is in



a good shape. We will talk about the details about the prediction model and the construction process in next chapter.

# Chapter 5

## Predictive Model

In this chapter, we are going to present the predictive model which is a key part of our design. In the previous chapter, we have known that there are two major phases in our  $B^p$ -tree, both of which will be based on this predictive model. Thus in this chapter, we will talk about how we construct the predictive model and use it in these two phases. We also know that the accuracy of the model is critical to the performance of the whole indexing and thus we will present our strategy to evaluate the realtime status of the predictive model and make adjustment to make it work properly. There will be three major parts of this chapter including the predictive model for warm-up phase, predictive model for update phase and evaluating  $B^p$ -tree.

### 5.1 Predictive Model for Warm-up

In this section, we introduce a predictive model to construct a  $B^p$ -tree structure in the warm-up phase. We will present running examples to show what the predictive model is and how it is integrated into our  $B^p$ -tree. We first discuss how to

predict a  $B^p$ -tree skeleton (Section 5.1.1), and then propose to construct a  $B^p$ -tree (Section 5.1.2).

### 5.1.1 Predicting the $B^p$ -tree Skeleton

Suppose there are  $N$  keys in the DRAM  $B^+$ -tree, the height of the  $B^+$ -tree is  $h$ , and the branching factor is  $2m$ . The  $B^p$ -tree has the same height as the  $B^+$ -tree, but with a larger branching order  $2M = K * 2m$ , where  $K$  is an integer and  $K \geq 1$ .  $K$  can be set by the administrator. We can also predict  $K$  as follows.

Let  $T$  denote the total number of possible keys in the whole dataset. We estimate  $T$  using the numbers of keys in the histogram. Suppose the maximal number of keys in a bucket is  $A$  and the bucket width is  $W = \frac{U-L}{|B|}$ , thus there are at most  $W$  keys in a bucket. We use  $N \times \frac{W}{A}$  to predict the possible key number  $T$ . As there are  $T$  keys in  $B^p$ -tree and  $N$  keys in  $B^+$ -tree, we set  $K = \log_h \frac{T}{N}$ .

Obviously, if we overestimate  $K$ , the  $B^p$ -tree will turn out to be sparse; on the contrary, if we underestimate the number, we may need to do more splits. We assume that each tree node in the final  $B^p$ -tree is expected to be  $\mu\%$  full, i.e., each leaf node has  $E = \mu\% \times 2M$  keys. Thus the  $i$ -th level is expected to have  $E^i$  nodes (the root is the first level, which has only one node).

After we got the number  $K$ , we can build the skeleton of the  $B^p$ -tree, in other words, we are going to enlarge the size of each node in the tree by  $K$ , in which case, there may be many “underflow” nodes. But we will not worry about this and we will gradually insert keys into these nodes and raise the average node utilization.

### 5.1.2 $B^p$ -tree Construction

In this section, we discuss how to construct the  $B^p$ -tree based on the current  $B^+$ -tree structure. We traverse the  $B^+$ -tree in post-order. For each node, we predict whether we need to split it based on our predictive model (which will be introduced later). If we need to split it, we split it into two (or more) nodes, and insert separators (or keys) into its parent node, which may in turn cause the parent node to be split. Following this approach, the process will continue to spread to the whole tree. As we employ a post-order traversal, we can guarantee that the child splits are before the parent split, and our method can keep a balanced tree structure.

Next we discuss how to split a  $B^+$ -tree node. For ease of presentation, we first introduce a concept.

**Definition 5.1 (Node Extent)** *Each node  $n$  in the index tree is associated with a key range  $[n_l, n_u]$ , where  $n_l$  and  $n_u$  are respectively the minimum key and the maximum key that could fall in this node. We call this range the extent of the node.*

The extent of node  $A$  and  $B$  of the  $B^+$ -tree in Figure 4.2, for example, are  $[0, 19)$  and  $[19, 32)$  respectively. If a node is not the leftmost or the rightmost child of its parent, we can get its extent from its parent (except for the root node); otherwise we need to determine it from its ancestors. In practice, we need to update this information each time we split or merge nodes and we also need to consider this overhead in the evaluation.

Consider a  $B^+$ -tree node  $n$  on the  $i$ -th level. Suppose its extent is  $[\text{key}_{\min}, \text{key}_{\max}]$  and currently it has  $|n|$  keys,  $\text{key}_1, \text{key}_2, \dots, \text{key}_{|n|}$ . We access the keys in order. Suppose the current key is  $\text{key}_j$ . We next discuss whether to split the node ac-

cording to  $\text{key}_j$  as follows. As  $\text{key}_{\min}$  is the possible minimum key in the node, we first estimate the number of possible keys between  $\text{key}_{\min}$  and  $\text{key}_j$ , denoted by  $P(\text{key}_{\min}, \text{key}_j)$ . Then we estimate the number of keys that can be accommodated between  $\text{key}_{\min}$  and  $\text{key}_j$  on the  $B^p$ -tree, denoted by  $A(\text{key}_{\min}, \text{key}_j)$ .

Obviously if  $P(\text{key}_{\min}, \text{key}_j) < A(\text{key}_{\min}, \text{key}_j)$ , we do not need to split the node according to  $\text{key}_j$ ; otherwise we need to split the node. We generate a  $B^p$ -tree node with keys  $\text{key}_{\min}, \dots, \text{key}_{j-1}$ , remove the keys  $\text{key}_{\min}, \dots, \text{key}_{j-1}$  from the DRAM  $B^+$ -tree node, insert the key  $\text{key}_j$  to its parent on DRAM  $B^+$ -tree, and update the pointers of the  $B^+$ -tree node: the left pointer of this key points to the  $B^p$ -tree node, and the right pointers of this key points to the  $B^+$ -tree node. Next we repeatedly split the node with keys  $\text{key}_j, \dots, \text{key}_{|n|}$  (Note that  $\text{key}_j$  turns to the first key in the new node). If we cannot split the node for the last key, we will create a  $B^p$ -tree node with the same keys in the  $B^+$ -tree node, and update the pointer of its parent to the  $B^p$ -tree node. Next we discuss how to predict  $A(\text{key}_{\min}, \text{key}_j)$  and  $P(\text{key}_{\min}, \text{key}_j)$ .

### Predicting the number of possible keys between $\text{key}_{\min}$ and $\text{key}_j$ :

If  $\text{key}_{\min}$  and  $\text{key}_j$  are in the same bucket  $B_s$ , we can estimate  $P(\text{key}_{\min}, \text{key}_j)$  as follows. Based on the histogram, there are  $n_s$  keys in the bucket. Then the number of keys between  $\text{key}_{\min}$  and  $\text{key}_j$  can be estimated by  $(\text{key}_j - \text{key}_{\min}) \times \frac{n_s}{W}$ , where  $W$  is the bucket width. Thus the number of possible keys in the range is

$$P(\text{key}_{\min}, \text{key}_j) = K \times (\text{key}_j - \text{key}_{\min}) \times \frac{n_s}{W}, \quad (5.1)$$

if  $\text{key}_{\min}$  and  $\text{key}_j$  are in the same bucket  $B_s$ .

On the contrary, if  $\text{key}_{\min}$  and  $\text{key}_j$  are in different buckets, we estimate the

number as follows. Without loss of generality, suppose  $\text{key}_{\min}$  is in bucket  $B_s$  and  $\text{key}_j$  is in bucket  $B_e$ . Let  $B_s^u$  denote the upper bound of keys in bucket  $B_s$  and  $B_e^l$  denote the lower bound of keys in bucket  $B_e$ . Thus the number of keys between  $\text{key}_{\min}$  and  $\text{key}_j$  in bucket  $B_s$  is  $(B_s^u - \text{key}_{\min}) \times \frac{n_s}{W}$ . The number of keys between  $\text{key}_{\min}$  and  $\text{key}_j$  in bucket  $B_e$  is  $(\text{key}_j - B_e^l) \times \frac{n_b}{W}$ . Thus the total number of keys between  $\text{key}_{\min}$  and  $\text{key}_j$  is  $(B_s^u - \text{key}_{\min}) \times \frac{n_s}{W} + \sum_{t=s+1}^{e-1} n_t + (\text{key}_j - B_e^l) \times \frac{n_b}{W}$ . Thus the number of possible keys between  $\text{key}_{\min}$  and  $\text{key}_j$  is

$$P(\text{key}_{\min}, \text{key}_j) = K \times \left( (B_s^u - \text{key}_{\min}) \times \frac{n_s}{W} + \sum_{t=s+1}^{e-1} n_t + (\text{key}_j - B_e^l) \times \frac{n_b}{W} \right), \quad (5.2)$$

if  $\text{key}_{\min}$  and  $\text{key}_j$  are in different buckets.

**Predicting the number of keys that can be accommodated between  $\text{key}_{\min}$  and  $\text{key}_j$ :**

Note that node  $n$  has  $|n|$  keys and it is expected to have  $E$  keys, thus the number of accommodated keys in this node is  $E - |n|$ . Thus

$$A(\text{key}_{\min}, \text{key}_j) = \min(\text{key}_j - \text{key}_{\min}, E - |n|),$$

if  $n$  is a leaf node.

If node  $n$  is a non-leaf node, we can directly add  $j$  children between the two keys. In addition, we can also add some keys between  $\text{key}_{\min}$  and  $\text{key}_j$  as there are  $E - |n|$  positions which are not used in the node. Obviously, we insert at most  $\min(\text{key}_j - \text{key}_{\min}, E - |n|)$  keys in the node. Thus we can add at most  $c = j + \min(\text{key}_j - \text{key}_{\min}, E - |n|)$  children under the node between  $\text{key}_{\min}$  and

$\text{key}_j$ . As node  $n$  is in the  $i$ -level, the children are on  $i + 1$ -level. As each child can have  $E$  keys and  $E + 1$  children, each node can have  $(E + 1)^{h-i-1}$  descendants. Thus there are  $c \times \sum_{t=0}^{h-i-1} (E + 1)^t$  nodes between the two keys. As each node can accommodate  $E$  keys, the total number of accommodated keys is

$$A(\text{key}_{\min}, \text{key}_j) = E \times c \times \sum_{t=0}^{h-i-1} (E + 1)^t, \quad (5.3)$$

if node  $n$  is a non-leaf node.

To summarize, we can use the predicted numbers to split the nodes. Iteratively, we can split all the nodes and insert the new nodes into PCM. Figure 5.1 illustrates the algorithm. The **Warm-up** algorithm first traverses the **B<sup>+</sup>-tree** in post-order by calling its function **PostOrder** (line 3). Function **PostOrder** splits the nodes iteratively. Given a node  $n$  on level  $i$ , it checks whether the node should be split by calling function **Split** (line 4), which is used to split a node based on our predictive model. If our model decides to split node  $n$ , we generate a **B<sup>p</sup>-tree** node with keys,  $\text{key}_{\min}, \dots, \text{key}_{j-1}$  (line 6), remove the keys,  $\text{key}_{\min}, \dots, \text{key}_{j-1}$ , from the DRAM **B<sup>+</sup>-tree** node (line 7), insert the key  $\text{key}_j$  to its parent on DRAM **B<sup>+</sup>-tree** (line 8), and update the pointers of the **B<sup>+</sup>-tree** node: the left pointer of this key to the **B<sup>p</sup>-tree** node, and the right pointers of this key to the **B<sup>+</sup>-tree** node. Next we repeatedly split the node with keys,  $\text{key}_j, \dots, \text{key}_{|n|}$  (line 9). If we cannot split the node for the last key, we will create a **B<sup>p</sup>-tree** node with the keys, and update the pointer of its parent to the **B<sup>p</sup>-tree** node (line 10). Iteratively, we can construct the **B<sup>p</sup>-tree** structure.

We show an example of the split algorithms in Figure 4.2. Here we take the leaf node for an example instead of the whole split algorithm. As can be seen from the

<b>Algorithm 1:</b> Warm-up( $B^+$ -tree, Histogram)	
<b>Input:</b> $B^+$ -tree and Histogram in DRAM Buffer	
<b>Output:</b> $B^p$ -tree on PCM	
1	<b>begin</b>
2	Let $r$ denote the root of the $B^+$ -tree, Level $i = 0$ ;
3	PostOrder ( $r, i$ , Histogram) ;
4	<b>end</b>
<hr/>	
<b>Function</b> PostOrder( $n, i$ , Histogram)	
<b>Input:</b> $n$ : $B^+$ -tree node; $i$ : Level of $n$ ; Histogram	
<b>Output:</b> $B^p$ -tree nodes	
1	<b>begin</b>
2	<b>for</b> each child $c$ of $n$ <b>do</b>
3	PostOrder ( $c, i + 1$ , Histogram);
4	$key_j = \text{Split}(n, i, \text{Histogram})$ ;
5	<b>while</b> $key_j \neq \phi$ <b>do</b>
6	Generate a $B^p$ -tree node with $key_{min}, \dots, key_{j-1}$ ;
7	Remove keys after $key_j$ from $B^+$ -tree node $n$ ;
8	Insert $key_j$ to the parent of $n$ on $B^+$ -tree and update the pointers ;
9	Split( $n, i$ , Histogram);
10	Create a node with $key_{min}, \dots, key_{j-1}$ , remove $n$ , update $n$ 's parent;
11	<b>end</b>
<hr/>	
<b>Function</b> Split( $n, i$ , Histogram)	
<b>Input:</b> $n$ : $B^+$ -tree node; $i$ : Level of $n$ ; Histogram	
<b>Output:</b> Key: Split Key	
1	<b>begin</b>
2	Let $key_{min}$ denote the first key in node $n$ ;
3	<b>for</b> $j = 2, 3, \dots  n $ <b>do</b>
4	<b>if</b> PredictTwoKeys ( $key_{min}, key_j, i$ ) <b>then</b>
5	return $key_j$ ;
6	return $\phi$ ;
7	<b>end</b>
<hr/>	
<b>Function</b> PredictTwoKeys( $key_{min}, key_j$ , Histogram)	
<b>Input:</b> $key_{min}$ ; $key_j$ ; Histogram	
<b>Output:</b> True or false	
1	<b>begin</b>
2	Compute $A(key_{min}, key_j)$ ; Compute $P(key_{min}, key_j)$ ;
3	<b>if</b> $A(key_{min}, key_j) < P(key_{min}, key_j)$ <b>then</b> return true;
4	<b>else</b> return false;
5	<b>end</b>

Figure 5.1: Warmup Algorithm



Figure 4.2, the previous node  $C$  is split into two nodes  $C$  and  $C'$ , though it is not full. Since there are only two keys in previous node  $C$ , we shall calculate  $A(\text{key}_{min}, \text{key}_2)$  and  $P(\text{key}_{min}, \text{key}_2)$  as follows.  $A(\text{key}_{min}, \text{key}_2) = \min(39 - 32, 4 - 2) = 2$ ,  $P(\text{key}_{min}, \text{key}_2) = (39 - 32) * \frac{8}{20} = 2.8$ . As  $A(\text{key}_{min}, \text{key}_2) < P(\text{key}_{min}, \text{key}_2)$ , according to the algorithms in Figure 5.1, node  $C$  needs to be split and a new node  $C'$  is created and then we update the pointers.

## 5.2 Predictive Model for Updates

In this section we propose a predictive model for the update phase. We will describe the basic operations on a  $B^p$ -tree, including search (Section 5.2.1), deletion (Section 5.2.2), and insertion (Section 5.2.3). Actually the search and deletion processes are similar to that of the standard  $B^+$ -tree (except the need to deal with unordered leaf entries). Thus our main focus will be on the insertion.

### 5.2.1 Search

Since we use a small DRAM as a buffer, some newly inserted keys will still be in the main memory  $B^+$ -tree and have not been merged into the main  $B^p$ -tree on PCM. Thus, besides searching the main  $B^p$ -tree in the manner similar to that of the  $B^+$ -tree, for our  $B^p$ -tree, a search operation still needs to search the small  $B^+$ -tree first. Then there will be two steps in the search process. We first lookup the small  $B^+$ -tree in the buffer, and then search the main  $B^p$ -tree. As noted, since the entries within a leaf node of the  $B^p$ -tree may not be sorted, the search will examine every key entry. If neither of the two steps return any results, **null** will be returned. The above steps are summarized in Figure 5.2. Obviously the time complexity of the

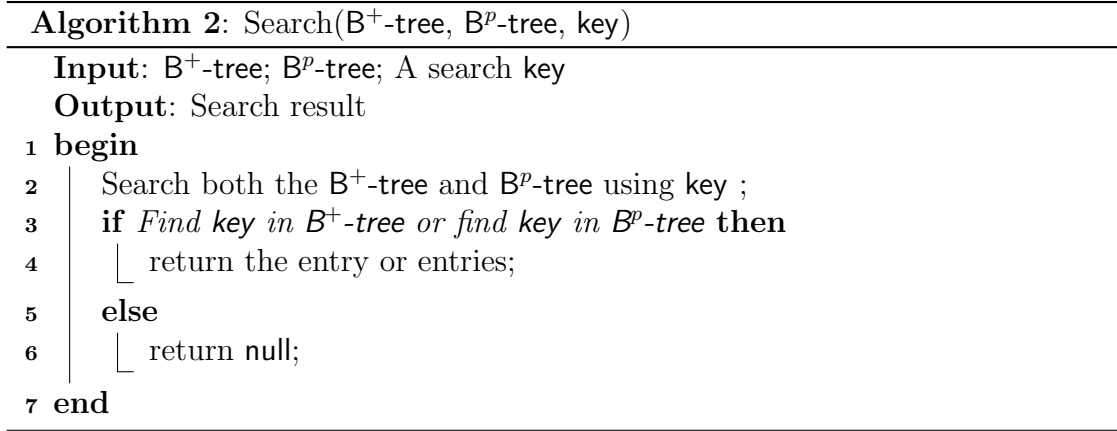


Figure 5.2:  $B^p$ -tree: Search Operation

search operation is  $\mathcal{O}(h)$ , where  $h$  is the height of the  $B^p$ -tree and it is similar to the traditional  $B^+$ -tree.

### 5.2.2 Deletion

Like search, deletion also requires searching both  $B^p$ -tree and  $B^+$ -tree. A deletion on the  $B^p$ -tree is handled in a similar way as that for standard  $B^+$ -tree, but with some differences. First, the deleted entry can be replaced by the last key entry in the node. This is to pack the entries within the leaf node. Second, if the corresponding leaf node has fewer than  $M$  keys, we will not borrow keys from its siblings. This can avoid the merge operations. The reason is that since the read latency of PCM is much shorter than the write latency, the overhead caused by the empty node in the query processing stage is negligible. Furthermore, the space could be reserved for the future keys to reduce subsequent writes.

Given a key to delete, we first search it from the  $B^+$ -tree. If we find the entry, we directly remove it from  $B^+$ -tree. If not, we then search it in the  $B^p$ -tree. If we find the leaf node in the  $B^p$ -tree, we remove the key from the node. Note that we will

not do merge operations even if the node has less than half ( $M$ ) keys. We do not propagate the deletion operation to its ancestors. The above steps are summarized in Figure 5.3. Obviously the time complexity of the deletion operation is  $\mathcal{O}(h)$ .

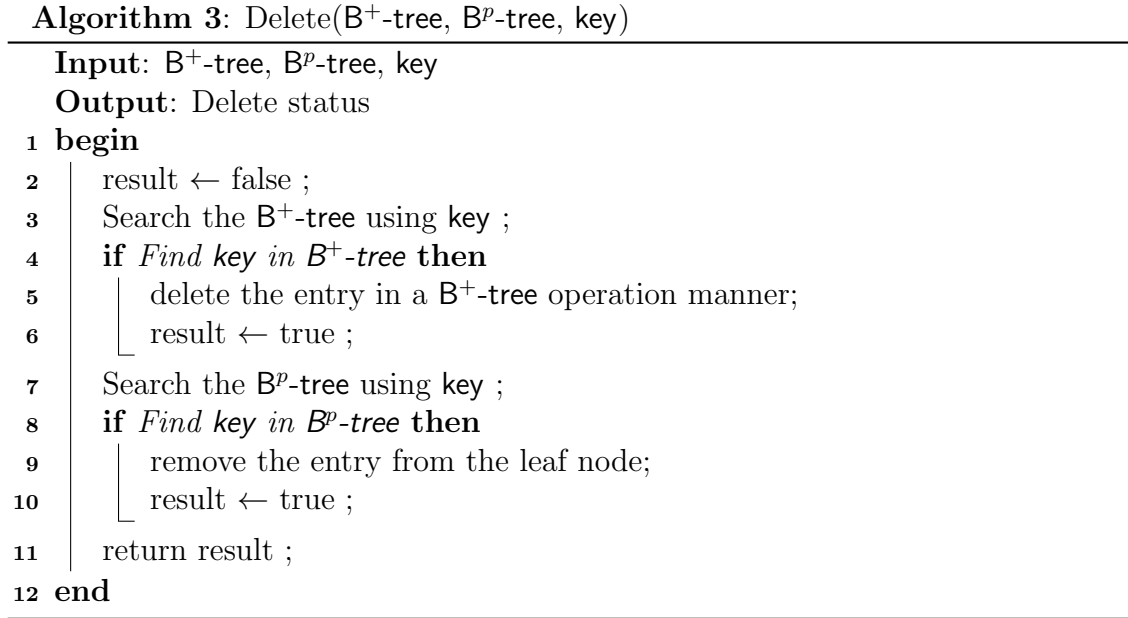


Figure 5.3:  $B^p$ -tree: Deletion Operation

### 5.2.3 Insertion

Since  $B^p$ -tree is maintained with the aid of a DRAM buffer and a predictive model, both the  $B^+$ -tree in the buffer and the histogram of the predictive model need to be updated in each insertion. When the buffer is full, the  $B^+$ -tree will be merged into the main  $B^p$ -tree on PCM.

All the keys in the  $B^+$ -tree will be inserted into the main tree one by one. Once a key is to be inserted, we first look up the leaf node  $L$  that the new key belongs to as the standard  $B^+$ -tree. Then we predict whether it should be directly inserted into the node or the node should be split. We first compute the number of keys

that can be accommodated in this node  $L$ , denoted by  $ANO_L$ . We then predict the number of keys that could fall in this node, denoted by  $PNO_L$ . If  $PNO_L \geq ANO_L$ , we need to split the node; otherwise, we will not. If we need to split the node, a new leaf node will be created and a “middle” key will be chosen based on the predictive model and pushed upward to the parent node. Existing keys in the node  $L$  needs to be adjusted according to the “middle” key.

Note that the middle key is not the key in the median position as the standard  $B^+$ -tree. Instead, we need to select a median key based on the data distribution (which will be discussed later). As we insert a middle key into its parent, it may cause its parent to split. The above steps are summarized in Figure 5.4. Next, we discuss how to compute  $PNO_n$  and  $ANO_n$  for node  $n$ .

#### **Computing the accommodated key number of node $n$ , $ANO_n$ :**

Suppose node  $n$  is in the  $i$ -th level. Each node has at most  $2M$  keys and  $2M+1$  pointers, thus node  $n$  has  $\sum_{t=1}^{h-i} (2M+1)^t$  descendants. Thus the accommodated key number of node  $n$  is

$$ANO_n = 2M * \sum_{t=0}^{h-i} (2M+1)^t. \quad (5.4)$$

#### **Predicting the possible key number occupancy in node $n$ , $PNO_n$ :**

Next we predict the total number of keys that could potentially belong to this node. We first find the extent of this node, denoted by  $[\text{key}_{min}, \text{key}_{max}]$ , where  $\text{key}_{min}$  and

$\text{key}_{max}$  are respectively the minimum key and the maximum key in this node. Based on the two bounds, we can compute the number of possible keys fell into this node as discussed in Section 5.1.2.

That is if  $\text{key}_{max}$  and  $\text{key}_{min}$  are in the same bucket  $B_s$ ,

$$PNO_n = K \times (\text{key}_{max} - \text{key}_{min}) \times \frac{n_s}{W}; \quad (5.5)$$

otherwise if  $\text{key}_{min}$  and  $\text{key}_{max}$  are respectively in two different buckets  $B_s$  and  $B_e$ .

$$PNO_n = K \times ((B_s^u - \text{key}_{min}) \times \frac{n_s}{W} + \sum_{t=s+1}^{e-1} n_t + (\text{key}_{max} - B_e^l) \times \frac{n_b}{W}). \quad (5.6)$$

Based on  $ANO_n$  and  $PNO_n$ , we can decide whether to split a node  $n$ . Next we discuss how to select a middle key if we need to split a node.

### Computing the middle key in node $n$ , $midKey$ :

Consider the keys in  $n$  are  $\text{key}_1, \text{key}_2, \dots, \text{key}_{|n|}$ . Without loss of generality, suppose  $\text{key}_1 \leq \text{key}_2 \leq \dots \leq \text{key}_{|n|}$ . Based on extent of a node, we define the middle key formally.

**Definition 5.2 (Middle Key)** *A key  $\text{key}_i$  in node  $n$  is called a middle key if*

$$P(\text{key}_{min}, \text{key}_i) \leq \frac{P(\text{key}_{min}, \text{key}_{max})}{2},$$

$$P(\text{key}_{min}, \text{key}_{i+1}) > \frac{P(\text{key}_{min}, \text{key}_{max})}{2},$$

where  $P(\text{key}_i, \text{key}_j)$  denote the number of predicted keys between  $\text{key}_i$  and  $\text{key}_j$ .

A straightforward method to find the middle key from a node is to compute  $P(\text{key}_{\min}, \text{key}_i)$  for each  $i$  from 1 to  $|n|$  until we find the middle key. The complexity is  $\mathcal{O}(M)$ . If the keys are sorted, e.g., the keys in an internode, we can use an alternative method. We have an observation that if the keys are sorted,  $P(\text{key}_{\min}, \text{key}_i) \leq P(\text{key}_{\min}, \text{key}_j)$  for  $i < j$  as formalized in Lemma 5.1. Thus we can employ a binary search method to find the middle key and reduce the time complexity to  $\mathcal{O}(\log M)$ . If the keys are unsorted, the complexity is  $\mathcal{O}(M)$ .

**Lemma 5.1** *Given a node  $n$  with keys ordered as  $\text{key}_1, \text{key}_2, \dots, \text{key}_{|n|}$ , and two keys  $\text{key}_{\min} \leq \text{key}_1$  and  $\text{key}_{\max} \geq \text{key}_{|n|}$ , we have*

$$P(\text{key}_{\min}, \text{key}_i) \leq P(\text{key}_{\min}, \text{key}_j)$$

for  $i < j$ .

Thus the worst-case time complexity of an insertion operation is  $\mathcal{O}(M + h \times \log M)$ , where  $M$  is the branching factor and  $h$  is the height.

### 5.3 Evaluating $\mathbf{B}^p$ -tree

In this section, we introduce several metrics to evaluate the status of  $\mathbf{B}^p$ -tree and use them to guide the future prediction which is necessary to keep the tree balanced.

The first metric is insertion overflow. When inserting a new entry into the  $\mathbf{B}^p$ -tree, we employ a leaf-to-root way, that is we always insert a key into leaf node first. If the node overflows, it needs to be split and some of the keys need to be

<b>Algorithm 4:</b> Insert( $B^+$ -tree, $B^p$ -tree, key)	
<b>Input:</b> $B^+$ -tree, $B^p$ -tree, key	
<b>Output:</b> Updated $B^p$ -tree	
1	<b>begin</b>
2	Search the leaf node for key, $L$ ;
3	Upgrade ( $L$ , Histogram);
4	<b>end</b>
<hr/>	
<b>Function</b> Upgrade( $n$ , <i>Histogram</i> )	
<b>Input:</b> $n$ : A $B^p$ -tree node, Histogram	
1	<b>begin</b>
2	$PNO_n \leftarrow \text{GetPredictedNumber}(n)$ ;
3	$ANO_n \leftarrow \text{GetAccomodatedNumber}(n)$ ;
4	<b>if</b> $PNO_n < ANO_n$ <b>then</b>
5	insert key into $n$ ;
6	return;
7	<b>else</b>
8	$midKey \leftarrow \text{GetMiddleKey}(n)$ ;
9	$midKey$ is pushed upward to the parent $p$ ;
10	A new leaf node $n'$ is created and new pointer from the parent node to $n'$ ;
11	Remove keys larger than $midKey$ from $n$ to $n'$ ;
12	Upgrade ( $p$ , Histogram) ;
13	<b>end</b>

Figure 5.4:  $B^p$ -tree: Insertion Operation

moved to other nodes. Obviously, the larger the number of keys in a node is, the higher is the probability for it to split and the overhead incurred to move keys. Thus we can use the number of keys in a node to evaluate the degree of insertion overflow. Given a node  $n$  with  $n_{keys}$  keys. A larger  $n_{keys}$  implies there is a higher probability for  $n$  to split, resulting in a larger number of possible writes.

The second metric is unqualified-node ratio. A node is called an *unqualified node* if its key number is smaller than  $M$ . If there are many unqualified nodes, the constructed  $B^p$ -tree is very sparse. For a node  $n$ , the smaller the value of  $n_{keys}$ , the

sparser the tree will be. To evaluate the overall **B<sup>p</sup>-tree**, we need to consider all tree nodes. Let  $n_{un}$  denote the number of unqualified nodes. The larger the value of  $n_{un}$ , the sparser the **B<sup>p</sup>-tree**. We can easily determine  $n_{un}$  as follows. Initially, all nodes are unqualified nodes and  $n_{un}$  is the total number of nodes. When inserting a key, if an unqualified node turns to be a qualified node (with key number no smaller than  $M$ ), we decrease the number  $n_{un}$  by 1.

Next we combine the above two factors to evaluate a **B<sup>p</sup>-tree**. As the expected utilization is  $\mu\%$  and then the average key number of a node is  $\mu\% \times 2M$ , we can use the following equation to evaluate the **B<sup>p</sup>-tree**,

$$Q = \sum_n \delta \times (n_{keys} - \mu\% \times 2M), \quad (5.7)$$

where

$$\delta = \begin{cases} \frac{n_{keys}}{\text{key}_{max} - \text{key}_{min}} & n_{keys} \geq \mu\% \times 2M \\ \frac{\text{key}_{max} - \text{key}_{min}}{n_{keys}} & n_{keys} < \mu\% \times 2M \end{cases} \quad (5.8)$$

and  $[\text{key}_{min}, \text{key}_{max}]$  is the extent of node  $n$ , in [7]  $\mu$  is 69 for the standard **B<sup>+</sup>-tree**.

If  $Q$  is larger than 0, then it means that the **B<sup>p</sup>-tree** is very dense. The larger the value of  $Q$ , the denser the **B<sup>p</sup>-tree** will be. However it may involve many more numbers of writes when the tree needs to be reorganized (by splits and merges). If  $Q$  is larger than an upper bound  $\tau_u$ , we need to tune our model to do more (planned) splits (when merging **B<sup>+</sup>-tree** with **B<sup>p</sup>-tree**).

On the contrary, if  $Q$  is smaller than 0, **B<sup>p</sup>-tree** is very sparse. The smaller the value of  $Q$ , the sparser the **B<sup>p</sup>-tree**. If  $Q$  is smaller than a lower bound  $\tau_l$ , we need



to tune our method to reduce the number of (planned) splits.

## 5.4 Summary

In this chapter, we introduced the predictive model we use to construct the  $B^p$ -tree. There are two phases when constructing the tree including the warm-up phase and the update phase. The warm-up phase means the time before the first time that the DRAM buffer is full and after that we start the update phase. In both phases, the predictive model is very important and it will influence whether we will split a node in advance or not. Then we talked about the normal operations of our  $B^p$ -tree and we gave the detailed algorithms. After that we proposed a evaluating model to evaluate the realtime status of our predictive model and some metrics are adopted to make it work properly and further ensure a good status of the whole indexing tree.

# Chapter 6

## Experimental Evaluation

In this chapter, we evaluate the performance of our proposed  $B^p$ -tree and show the experimental results. An extensive performance study is conducted to show the efficiency and effectiveness of the  $B^p$ -tree with various types of queries and updates. Since we have not got any PCM prototype in hand, we need to setup an environment to simulate the PCM characteristics first. Then we did our experiments in the simulated environment. Our results show  $B^p$ -tree outperforms the traditional  $B^+$ -tree on the insertion and deletion performance, while holding a similar search performance at the same time.

### 6.1 Experimental Setup

In this section, we are going to present our experimental setup. We talk about the experimental platform first and we will present the details of our simulation platform. After that we will propose our data set, workloads and the different algorithms we are going to compare with.

### 6.1.1 Experimental platform

We integrate our proposed  $B^p$ -tree in PostgreSQL and extend the buffer management module to support the PCM model. We follow the specifications of the PCM model in [6] and we also talked about this in Chapter 2. Three metrics are used in our experiments to measure the performance the  $B^p$ -tree, namely the number of writes, energy consumption and CPU cycles. Each time when we write a new cache line into the PCM, we compute the number of modified bits and bytes by comparing it with the previous one. In our experiments, the number of writes is computed as the number of modified bytes, while the energy consumption is estimated by the number of modified bits. We compute the CPU cost by combining the CPU cycles of our  $B^p$ -tree in both PCM and DRAM.

The experiments were conducted in CentOS release 5.6 with g++ 4.1.2. Our system is powered with a 16-core Intel Xeon E5620 2.4GHz CPU and 64GB main memory. Based on the benchmark used in [6, 27, 5], we set the parameters as follows: the read latency of a PCM cache line is 288 cycles; the write latency of PCM is 281 cycles for every 4 bytes; the read latency of a DRAM cache line is 115 cycles and the write latency of DRAM is 7 cycles for every 4 bytes. In PCM, the energy consumption is estimated as: the read energy per bit is 2pJ and the write energy per bit is 16pJ. In Table 6.1, we list the other parameters used in our experiments and their value ranges.

### 6.1.2 Data sets and workloads

Two synthetic datasets are used in our experiments. One is generated to follow the uniform distribution, while the other one follows the skewed distribution. We

Table 6.1: Parameters and their value ranges

Parameter	Value Ranges
Size of DRAM buffer	5% of the size of the PCM used
Size of cache line	64B
Size of the $B^+$ -tree node	256B (4 cache lines)
Size of the $B^p$ -tree node	256B, 512B, 1024B (4, 8, 16 cache lines)
K	1, 2, 4
Number of keys in the data set	5 millions

generate 5 millions keys in each dataset. In our experiments, the node size of the DRAM  $B^+$ -tree is 256B, which is equivalent to 4 cache lines; whereas the node size of all tree structures on the PCM varies from 256B, 512B to 1024B. In our  $B^p$ -tree, each index entry contains a 4-Byte key and a 4-Byte pointer. The size of the DRAM buffer used is approximately 5% of the size of the PCM. We generate various workloads (i.e., insertions, updates and searches) to study the performance of our approach. Specifically, an update is processed as a deletion operation followed by an insertion operation, and the search queries are composed of both the point queries and range queries. Based on our experimental results, we find that the performance on uniform dataset is similar to that on skewed dataset. Therefore, we only report the results on skewed dataset.

### 6.1.3 Algorithms compared

We compare four different indexing structures including our  $B^p$ -tree, the traditional  $B^+$ -tree, the proposed Unsorted Leaf tree in [6] and our  $B^p$ -tree with sorted leaf nodes on the basis of the following measures, the number of writes during the insertions, the energy consumption, the CPU cycles (including the small  $B^+$ -tree in the DRAM for  $B^p$ -tree) during the insertions and searches, and the leaf nodes

utilization.

As the  $B^p$ -tree is a composite structure including the main  $B^p$ -tree on PCM and the small buffer  $B^+$ -tree on DRAM, we need to determine how to compute each performance metric first. In the performance evaluation, we only consider the PCM cost while calculating the number of writes and the energy consumption. As for the CPU cost, the CPU cycles occupied for manipulating the DRAM  $B^+$ -tree are recorded, which can represent a more accurate processing time. All the indexes are tuned and only the best results are reported.

In all figures presented in this section, “BP-tree” represents our  $B^p$ -tree; “B-tree” represents the traditional  $B^+$ -tree; “Unsorted” represents the proposed unsorted leaf tree in [6]; and “BP-minus” represents  $B^p$ -tree with sorted leaf nodes. The x-axis represents the node size of the corresponding tree, e.g., x-coordinates 4 indicates that the node size of the corresponding tree is 4 cache lines.

## 6.2 Results and Analysis

We did various of experiments and in this chapter, we will show the results including insertion, update, search and node utilization. At last, we also did some more experiments to show that our  $B^p$ -tree indexing can work well under different data distributions.

### 6.2.1 Insertion

We first evaluate the insertion performance of  $B^p$ -tree. We insert all the keys in the dataset back-to-back using the different indexing algorithms. Moreover, for each data set, we build the tree using three different node sizes, that is, 4, 8, 16 cache

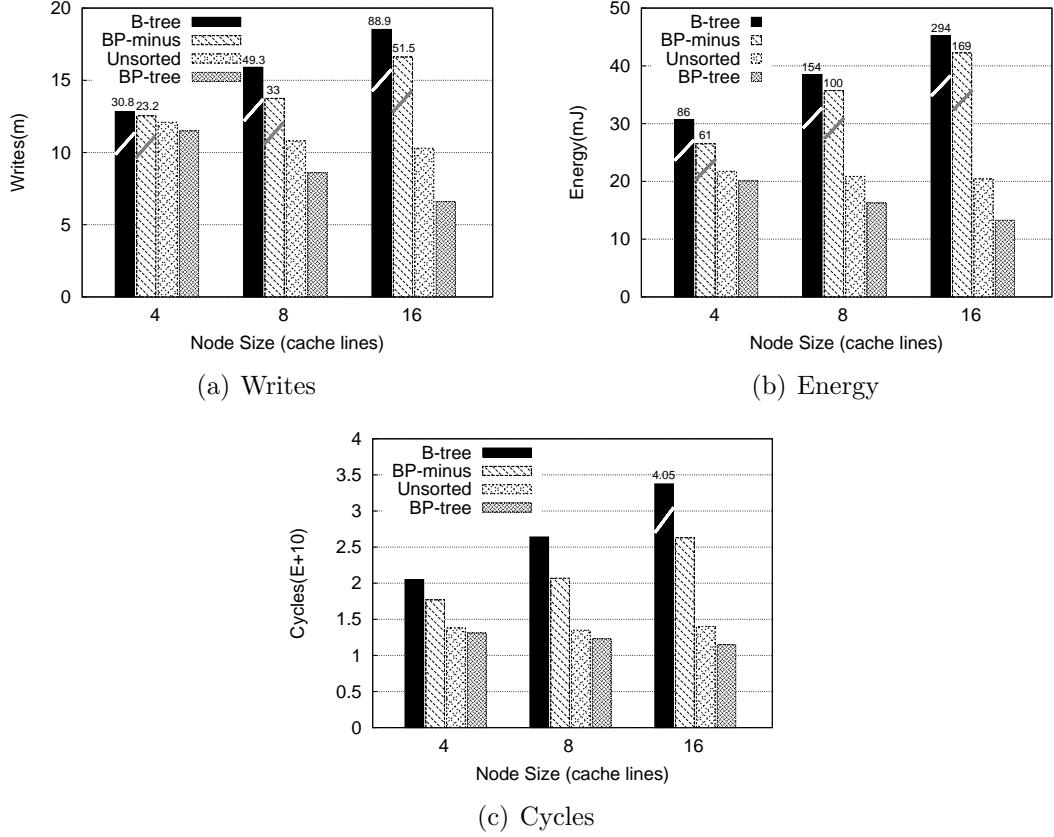


Figure 6.1: Insertion performance

lines, respectively.

In Figure 6.1, we compare the insertion performance of four tree indexing schemes. The three subfigures correspond to our three metrics respectively. In each subfigure, we present the performance of the four tree structures with three different node sizes. The scale of y-axis in Figure 6.1(a) and (b) are both in millions. We get two interesting observations from the results.

First, our  $B^p$ -tree achieves the best performance on all the three metrics and the performance gap increases as the node size becomes larger. The reason is that for large node sizes, our predictive model can estimate the splits more accurately, which can significantly reduce the number of writes by avoiding online splitting.

On the other hand, Unsorted outperforms  $B^+$ -tree and BP-minus. This is because most writes will appear in leaf nodes and Unsorted can reduce the number of writes on leaf nodes. Our  $B^p$ -tree outperforms the Unsorted scheme, as it splits the nodes in advance, which can reduce the numbers of future splits.  $B^p$ -tree incurs about 5%, 22%, 37% less PCM writes than the Unsorted scheme on the three different node sizes respectively. For energy consumption, the result is very similar to that of the writes. For CPU cycles, the gap becomes slightly smaller because  $B^p$ -tree incurs extra CPU costs on the small  $B^+$ -tree in the DRAM buffer. However,  $B^p$ -tree still performs better than the Unsorted tree by a factor of 18% when the node size is 16 cache lines.

Second, we compare the performance of the two tree indexes with sorted leaf nodes, namely BP-minus and  $B^+$ -tree. BP-minus outperforms  $B^+$ -tree in all metrics. BP-minus reduces about 25%, 33%, 42% of numbers of PCM writes compared to the  $B^+$ -tree. Similar trend is observed for the energy consumption. This means that our  $B^p$ -tree outperforms the traditional  $B^+$ -tree even if we do not want to make the keys on each node unsorted. For CPU cycles, the gap is not that significant because of the extra cost on the small  $B^+$ -tree in the DRAM buffer. Despite this, BP-minus still reduces 14%, 22%, 35% cost of that of  $B^+$ -tree.

### 6.2.2 Update

In this section, we evaluate the update performance of  $B^p$ -tree. We first insert all the keys back-to-back as the previous insertion experiment and then we generate and run 100k update queries randomly. The update query consists of two keys, *oldKey* and *newKey*. We first search the *oldKey*. If we find it, we delete it and insert the *newKey*. Otherwise, we will ignore the insertion request. In Figure 6.2,

we compare the average update performance of our  $B^p$ -tree and the other three tree structures. The result is very close to that of the insertion performance.

Our  $B^p$ -tree still achieves the best performance on all the three measures. The main reason is that our  $B^p$ -tree can predict future insertions and can pre-allocate space to reduce the number of writes. Compared to Unsorted, our  $B^p$ -tree reduces 24% of the writes, 26% of the energy and 19% of the CPU cycles, when the node size is 16 cache lines. If the node size is small, the gap decreases but our  $B^p$ -tree still outperforms Unsorted.

Compared to the traditional  $B^+$ -tree, the performance of BP-minus is better. It reduces 14% of the writes, 22% of the energy and 7% of the CPU cycles and the gap increases as the node size becomes larger. It shows the similar trends for all of the three measures.

### 6.2.3 Search

The philosophy of  $B^p$ -tree is two-fold: 1)  $B^p$ -tree is designed to reduce the number of writes on PCM, and 2)  $B^p$ -tree should be efficient for query processing as well. In this section, we evaluate the search performance of  $B^p$ -tree. The experiments include point queries and range queries. We experiment on both the uniform and skewed datasets. We first insert all keys into the index. Then for both point query and range query, we randomly generate 50k queries and calculate the CPU cycles during the processing.

In Figure 6.3, we compare the search performance of the four tree indexes. The left subfigure is for point query and the right one is for range query. The y-axis represents the total CPU cycles to run these search queries. For point query, the performance of  $B^p$ -tree is better than Unsorted. This is because when we process



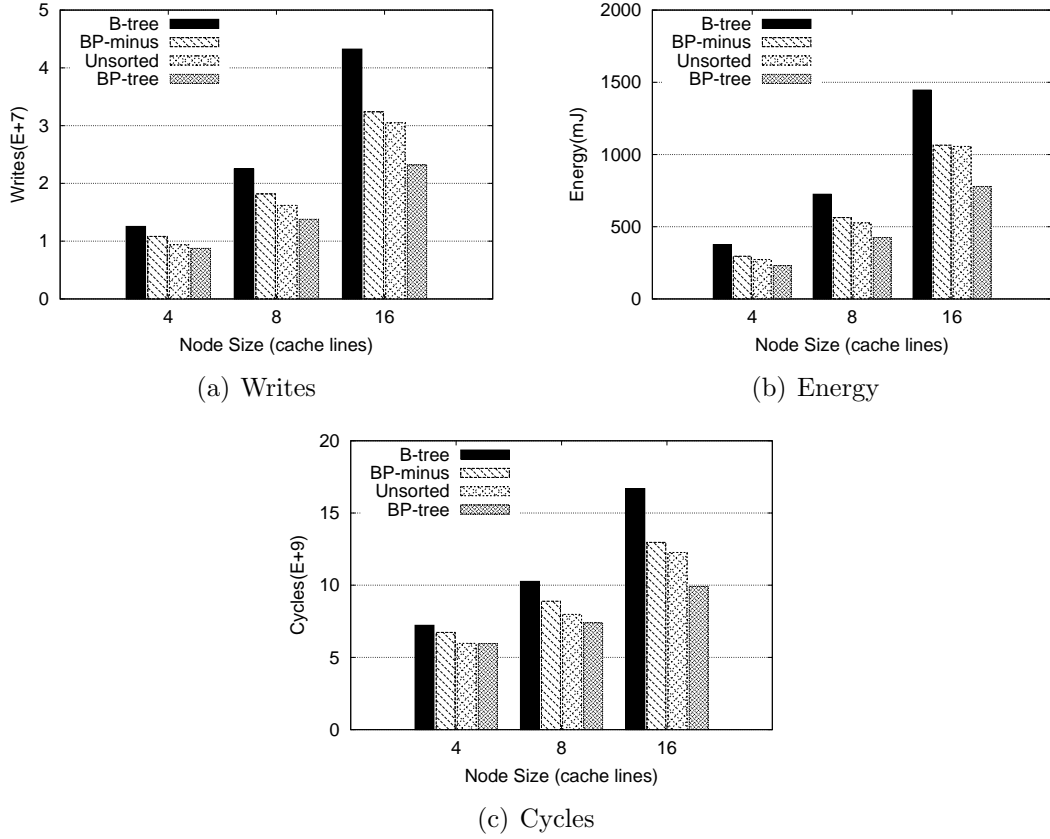


Figure 6.2: Update performance

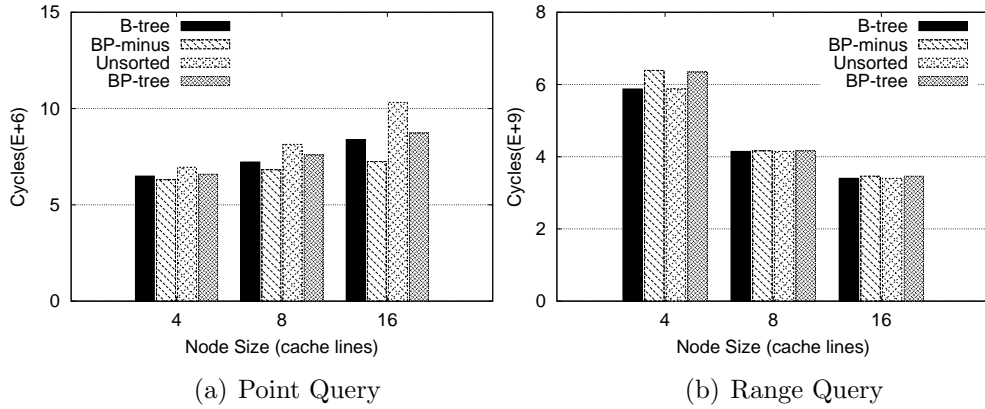


Figure 6.3: Search performance

the search query, we simply scan the node and once we find the key, we will return the associated data. If the  $B^p$ -tree has more leaf nodes than Unsorted, some keys

located in the right part of some nodes in Unsorted may be in the left part of some nodes in  $B^p$ -tree and thus more cache line reads are needed. The performance of BP-minus is better than that of Unsorted, which is expected since each search in Unsorted should read all the keys in the leaf node.

For range query, we can find that when the node size is 4 cache lines, the performance of our  $B^p$ -tree and BP-minus is worse than that of the  $B^+$ -tree and Unsorted. The reason is that when the node size is small, the tree will be more sensitive to the split strategy and generate more leaf nodes which could affect the range search performance. When the node size is larger, all the four tree indexes show a similar performance. This result is very important which means that the indexing tree is in a good shape and it did not split too “early” and make the tree too sparse.

#### 6.2.4 Node Utilization

In this experiment, we compare the leaf node utilization of the  $B^p$ -tree and the traditional  $B^+$ -tree. The experiments are same as the insertion performance experiment and we build the two trees based on the same data set and calculate the leaf nodes utilization periodically during the insertion. The scale of the x-axis is 0.5 million which means that 8 represents 4 millions keys inserted. The suffixes -4, -8, -16 in the figure indicate different node sizes.

As we can see in the figure, the leaf node utilization of the  $B^+$ -tree is stable, around 70% which is close to our assumption in Section 5.3. When the node size of the  $B^p$ -tree is 4 cache lines which is the same as that of the  $B^+$ -tree, the utilization is similar to that of the  $B^+$ -tree at first and then decreases as early splits happen and then it increases as the evaluation metrics described in Section 5.3 starts to




Figure 6.4: Leaf nodes utilization

work. When the node size is 8 cache lines, the utilization is smaller than that of the  $B^+$ -tree at first because of the enlargement of the node size and then it starts to increase. The result for the node size of 16 cache lines is similar. The stable utilization of all the three different  $B^p$ -tree indexes are all slightly smaller than that of the traditional  $B^+$ -tree, but according to the previous range search experiment, the influence of the utilization gap on the range search performance is not obvious and it is affordable.

### 6.2.5 Sensitivity to Data Distribution Changes

In this section, we evaluate the sensitivity of our predictive model to data distribution changes in order to show that our  $B^p$ -tree is stable with the dataset of different data distributions. We change the dataset as follows. The size of the dataset is 5 millions and the dataset follows a skewed (Zipf) distribution. However, we gradually change the Zipf factors and add a random offset every one million keys generated, resulting in a change of the data distribution. We did the insertion, update, search experiments as in previous sections. In Figure 6.5, we show comparisons of the

CPU cycles of all the four tree indexes with respect to different operations.

From the figure, we can observe that the relative performance of insertion, update and point search is very similar to that of the previous experiments. For the leaf node utilization, when the node size is 4 cache lines, the trend of the first half is similar to that of the previous result, but the utilization decreases slightly as the second half starts and increases again at last. The reason of the decrease is that changes of data distribution caused a wrong prediction from the predictive model and further caused some improper splits. After that the predictive model adjusts its prediction via the evaluation metrics and makes the structure normal again which means that our evaluating scheme works fine and it can help the predictive model to modify the splitting strategy.

We can also observe from Figure 6.5(e) that the stable utilization value is a bit smaller than that of the previous experiments, which may have also caused the range search performance to degrade slightly as shown in Figure 6.5(d). To summarize, the major performance of the  $B^p$ -tree verified in the previous experiments still holds when the data distribution changes which shows  $B^p$ -tree to be stable.

## 6.3 Summary

In this chapter, we did an extensive experiments evaluation of our  $B^p$ -tree. We build our experimental platform to simulate the PCM environment and compare it with some of the other write-optimized indexing technique and the traditional normal  $B^+$ -tree. We did the experiments based on both the uniform dataset and skewed dataset. We observed that for both data distribution, our  $B^p$ -tree can work well. The experimental results show that the  $B^p$ -tree significantly reduces the number of

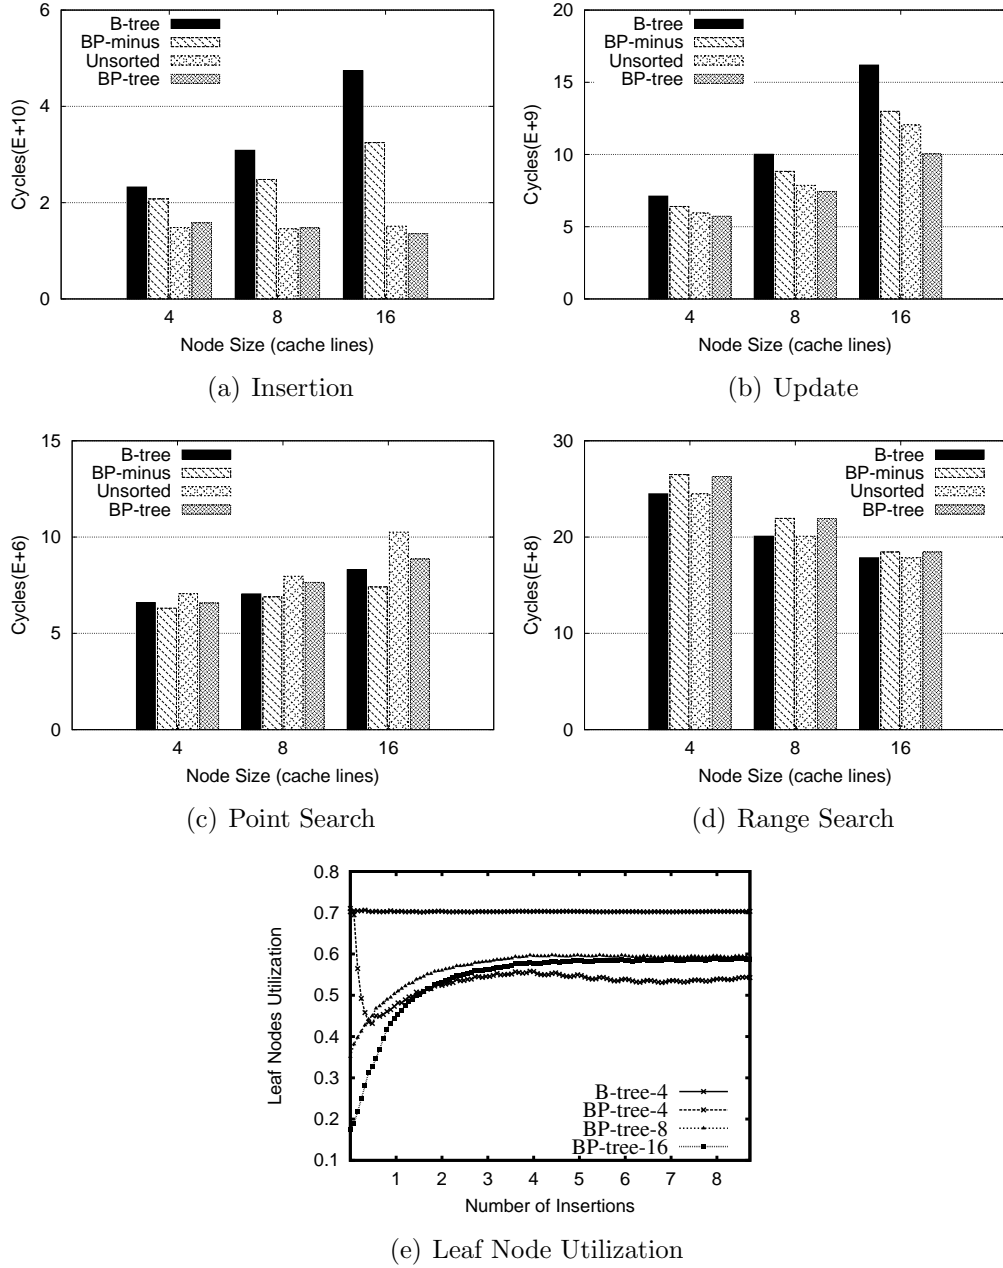


Figure 6.5: Sensitivity to Data Distribution Changes

writes for both the insertions and updates while having a good search performance at the same time, therefore making it write and energy efficient and suitable for a PCM-like hardware environment. The sensitivity experiments show that  $B^p$ -tree is

a stable indexing technique and it works fine for different datasets.

# Chapter 7

## Conclusion

In this chapter, we are going to conclude our work of this thesis and present a direction for the future work.

### 7.1 Conclusion

Phase change memory (PCM) is an emerging memory technology with many attractive features. In the near future, PCM is expected to become a common component in the memory hierarchy of the computer systems. On one hand, PCM could turn out to be a low-cost, more reliable, faster and better alternative to flash memory. On the other hand, PCM is also a very promising alternative to the traditional DRAM to become the major component of the main memory system.

However, PCM is still in its early stage and we still face some challenges to design algorithms for PCM-based memory systems including the high read and write latency compared to DRAM, the limited lifetime of the chip and the high energy consumption etc. If we want to make best use of PCM in the existing systems, we need to overcome these challenges.

In this thesis, we proposed to study the algorithms redesigning for database systems on phase change memory, particularly for indexing technique but in the future we can reconsider the whole technology stack of a traditional database systems. In our work, we proposed the  $B^p$ -tree indexing structure to best take advantage of the good features of PCM.

The main design objective of our  $B^p$ -tree indexing is to reduce the number of writes and energy consumption while keeping the tree construction and search efficient. We developed a predictive model to predict the near future data distribution based on the current data and then pre-allocate space for them in the PCM memory to reduce the movements caused by node splits and merges. We present the details of the model and show some metrics to evaluate the performance of the model during the construction process. If the metrics indicate that the prediction model is not running in a normal manner, we will adjust the model or rebuild the index in the worst case. The experiments on PostgreSQL database system showed that our  $B^p$ -tree indexing scheme achieves a better performance than the traditional  $B^+$ -tree and outperforms the state-of-the-art solution proposed in [6]. Additionally, our  $B^p$ -tree can be easily implemented in existing commercial database systems based on the existing  $B^+$ -tree structures.

## 7.2 Future Work

Database system is a very complex system and it consists of many complex subsystems[28, 9]. Currently we only considered the indexing technique but there are many others we can work on for example the query processing algorithms.

In the future, we can continue to work on algorithms redesigning for PCM-based



database systems. We can gradually move on the the query processing algorithms and buffer management. Since query processing is an important component of the database system, which will influence the overall performance of the whole system greatly, it needs to be redesigned carefully.

In [6], Chen, Gibbons and Nath have already done some work on hash joins. First they show the simple hash join, there are two phases, the build phase and the probe phase. In the build phase, the algorithm scans the smaller build relation and build a hash table for this relation. Then in the probe phase, the algorithm scans the larger probe relation. For each probe tuple, it computes the hash code and compare with the hash table to output the join result. The main disadvantage of the simple hash join is that whenever the size of the recode is large or not, it will lead to many cache misses, which can lead to a low performance. Then to solve the cache miss problem, the cache partitioning algorithm is introduced. The major difference is that in the build phase, instead of hashing the whole relation, the algorithm partitioned the two input relations using the same hash function so that every pair of the partitions can fit into the CPU cache. Then in the join phase, for each pair of the partition, the simple hash join algorithm can be used. The cache partitioning algorithm can solve the cache miss problem of the simple hash join. But it introduces a large number of writes which is not suitable for PCM-based database systems. Thus in this paper, the authors propose the virtual partitioning hash join, which is a variant of the cache partitioning method. The basic idea is that instead of physically copying input records into partitions, we perform the partitioning virtually. For each partition in build phase, we remember the record IDs and then in the join phase, we can use the record ID lists to join the records of a pair of partitions in place, thus avoiding the large number of writes in cache

partitioning.

In our research, since we want to reduce the writes, we need to reconsider and redesign the operators if necessary, including scan, sort, join etc. Since PCM does not have the “erase-before-write” shortage, we do not need to focus on the erase and block writes. Since PCM supports fast random reads, we can work on the page layouts like the algorithms design for SSD to reduce the read latency. For refinement of operators, some of the existing techniques can be used like the “first index and late materialization” strategy which can reduce much writes.

Much work can be done for the query processing optimization for PCM-based database systems. In this thesis, we just show a direction and we can continue to work on this topic in the future.

# Bibliography

- [1] D. Agrawal, D. Ganesan, R. Sitaraman, Y. Diao, and S. Singh. Lazy-adaptive tree: An optimized index structure for flash devices. *PVLDB*, 2(1):361–372, 2009.
- [2] Anastassia Ailamaki, David J. DeWitt, Mark D. Hill, and Marios Skounakis. Weaving relations for cache performance. In *Proceedings of the 27th International Conference on Very Large Data Bases*, VLDB '01, pages 169–180, San Francisco, CA, USA, 2001. Morgan Kaufmann Publishers Inc.
- [3] Lars Arge. The buffer tree: A new technique for optimal i/o-algorithms. *Algorithms and Data Structures*, pages 334–345, 1995.
- [4] D. Bausch, I. Petrov, and A. Buchmann. On the performance of database query processing algorithms on flash solid state disks. In *Database and Expert Systems Applications (DEXA), 2011 22nd International Workshop on*, pages 139–144. IEEE, 2011.
- [5] F. Bedeschi, R. Fackenthal, C. Resta, EM Donze, M. Jagasivamani, E. Buda, F. Pellizzer, D. Chow, A. Cabrini, G.M.A. Calvi, et al. A multi-level-cell bipolar-selected phase-change memory. In *Solid-State Circuits Conference*,

2008. *ISSCC 2008. Digest of Technical Papers. IEEE International*, pages 428–625. IEEE, 2008.
- [6] S. Chen, P.B. Gibbons, and S. Nath. Rethinking Database Algorithms for Phase Change Memory. CIDR, 2011.
- [7] Douglas Comer. The ubiquitous b-tree. *ACM Comput. Surv.*, 11(2):121–137, 1979.
- [8] Jeremy Condit, Edmund B. Nightingale, Christopher Frost, Engin Ipek, Benjamin Lee, Doug Burger, and Derrick Coetzee. Better i/o through byte-addressable, persistent memory. In *Proceedings of the ACM SIGOPS 22nd symposium on Operating systems principles, SOSP '09*, pages 133–146, New York, NY, USA, 2009. ACM.
- [9] CJ Date. An introduction to database systems , vol. ii addison-wesley. *Reading, MA*, 1983.
- [10] Goetz Graefe. Write-optimized b-trees. In *Proceedings of the Thirtieth international conference on Very large data bases-Volume 30*, pages 672–683, 2004.
- [11] Peter J. Haas and Joseph M. Hellerstein. Ripple joins for online aggregation. In *Proceedings of the 1999 ACM SIGMOD international conference on Management of data, SIGMOD '99*, pages 287–298, New York, NY, USA, 1999. ACM.
- [12] J.M. Hellerstein, M.J. Franklin, S. Chandrasekaran, A. Deshpande, K. Hildrum, S. Madden, V. Raman, and M.A. Shah. Adaptive query processing: Technology in evolution. *IEEE Data Engineering Bulletin*, 23(2):7–18, 2000.

- [13] Navin Kabra and David J. DeWitt. Efficient mid-query re-optimization of sub-optimal query execution plans. 27(2):106–117, June 1998.
- [14] Benjamin C. Lee, Engin Ipek, Onur Mutlu, and Doug Burger. Architecting phase change memory as a scalable dram alternative. In *Proceedings of the 36th annual international symposium on Computer architecture*, ISCA '09, pages 2–13. ACM, 2009.
- [15] Yinan Li, Bingsheng He, Robin Jun Yang, Qiong Luo, and Ke Yi. Tree indexing on solid state drives. *Proceedings of the VLDB Endowment*, 3(1):1195–1206, 2010.
- [16] Soumyadeb Mitra, Windsor W. Hsu, and Marianne Winslett. Trustworthy keyword search for regulatory-compliant record retention. In *Proceedings of the International Conference on Very Large Data Bases*, pages 1001–1012, 2006.
- [17] G. Muller, N. Nagel, C.U. Pinnow, and T. Rohr. Emerging non-volatile memory technologies. In *Proceedings of the 29th European Solid-State Circuits Conference*, ESSCIRC '03, pages 37–44. IEEE, 2003.
- [18] K.G. Murty. Supply chain management in the computer industry. Technical report, Working Paper, University of Michigan Department of Industrial and Operations Engineering, 2000.
- [19] K.G. Murty. Histogram, an ancient tool and the art of forecasting. Technical report, Citeseer, 2002.

- [20] Suman Nath and Aman Kansal. Flashdb: dynamic self-tuning database for nand flash. In *Proceedings of the 6th international conference on Information processing in sensor networks*, IPSN '07, pages 410–419. ACM, 2007.
- [21] Patrick O’Neil, Edward Cheng, Dieter Gawlick, and Elizabeth O’Neil. The log-structured merge-tree (lsm-tree). *Acta Informatica*, 33(4):351–385, 1996.
- [22] P.E. O’Neil. Thesb-tree an index-sequential structure for high-performance sequential access. *Acta Informatica*, 29(3):241–265, 1992.
- [23] Hwee Hwa Pang, Michael J. Carey, and Miron Livny. Partially preemptible hash joins. In *Proceedings of the 1993 ACM SIGMOD international conference on Management of data*, SIGMOD '93, pages 59–68, New York, NY, USA, 1993. ACM.
- [24] Hweehwa Pang, Michael J. Carey, and Miron Livny. Memory-adaptive external sorting. In *19th International Conference on Very Large Data Bases*, pages 618–629. Morgan Kaufmann, 1993.
- [25] Jian Pei, Man Ki Mag Lau, and Philip S. Yu. Ts-trees: A non-alterable search tree index for trustworthy databases on write-once-read-many (worm) storage. In *AINA*, pages 54–61, 2007.
- [26] Moinuddin K. Qureshi, John Karidis, Michele Franceschini, Vijayalakshmi Srinivasan, Luis Lastras, and Bulent Abali. Enhancing lifetime and security of pcm-based main memory with start-gap wear leveling. In *Proceedings of the 42nd Annual IEEE/ACM International Symposium on Microarchitecture*, MICRO '42, pages 14–23. ACM, 2009.

- [27] Moinuddin K. Qureshi, Vijayalakshmi Srinivasan, and Jude A. Rivers. Scalable high performance main memory system using phase-change memory technology. In *Proceedings of the 36th annual international symposium on Computer architecture*, ISCA '09, pages 24–33. ACM, 2009.
- [28] R. Ramakrishnan, J. Gehrke, and J. Gehrke. *Database management systems*. McGraw-Hill, 2003.
- [29] Vijayshankar Raman, Bhaskaran Raman, and Joseph M. Hellerstein. Online dynamic reordering. *The VLDB Journal*, 9(3):247–260, December 2000.
- [30] Luiz E. Ramos, Eugene Gorbatoov, and Ricardo Bianchini. Page placement in hybrid memory systems. In *Proceedings of the international conference on Supercomputing*, ICS '11, pages 85–95. ACM, 2011.
- [31] S. Raoux, GW Burr, MJ Breitwisch, CT Rettner, Y.C. Chen, RM Shelby, M. Salinga, D. Krebs, S.H. Chen, H.L. Lung, et al. Phase-change random access memory: A scalable technology. *IBM Journal of Research and Development*, 52(4.5):465–479, 2008.
- [32] Mendel Rosenblum and John K. Ousterhout. The design and implementation of a log-structured file system. In *Proceedings of the thirteenth ACM symposium on Operating systems principles*, SOSP '91, pages 1–15. ACM, 1991.
- [33] Stuart Schechter, Gabriel H. Loh, Karin Straus, and Doug Burger. Use ecp, not ecc, for hard failures in resistive memories. In *Proceedings of the 37th annual international symposium on Computer architecture*, ISCA '10, pages 141–152. ACM, 2010.

- [34] Nak Hee Seong, Dong Hyuk Woo, and Hsien-Hsin S. Lee. Security refresh: Prevent malicious wear-out and increase durability for phase-change memory with dynamically randomized address mapping. In *Proceedings of the 37th annual international symposium on Computer architecture*, ISCA '10, pages 383–394. ACM, 2010.
- [35] M.A. Shah, S. Harizopoulos, J.L. Wiener, and G. Graefe. Fast scans and joins using flash drives. In *Proceedings of the 4th international workshop on Data management on new hardware*, pages 17–24. ACM, 2008.
- [36] D. Tsirogiannis, S. Harizopoulos, M.A. Shah, J.L. Wiener, and G. Graefe. Query processing techniques for solid state drives. In *Proceedings of the 35th SIGMOD international conference on Management of data*, pages 59–72. ACM, 2009.
- [37] T. Urhan and M.J. Franklin. Xjoin: A reactively-scheduled pipelined join operator. *Bulletin of the Technical Committee on*, page 27, 2000.
- [38] H. Wong, S. Raoux, S.B. Kim, J. Liang, J.P. Reifenberg, B. Rajendran, M. Asheghi, K.E. Goodson, et al. Phase change memory. *Proceedings of the IEEE*, 98(12):2201–2227, 2010.
- [39] Chin-Hsien Wu, Tei-Wei Kuo, and Li Ping Chang. An efficient b-tree layer implementation for flash-memory storage systems. *ACM Transactions on Embedded Computing Systems (TECS)*, 6(3):19–es, 2007.
- [40] Weiye Zhang and Per-AAke Larson. Dynamic memory adjustment for external mergesort. In *Proceedings of the 23rd International Conference on Very Large*



*Data Bases*, VLDB '97, pages 376–385, San Francisco, CA, USA, 1997. Morgan Kaufmann Publishers Inc.

- [41] Ping Zhou, Bo Zhao, Jun Yang, and Youtao Zhang. A durable and energy efficient main memory using phase change memory technology. In *Proceedings of the 36th annual international symposium on Computer architecture*, ISCA '09, pages 14–23. ACM, 2009.